

Initial analysis

pablo paulsen

15/11/2021

```
data = read.csv('../Ftf Efficiency Dataset/ftf_ev_efficiency_distance_models_20211006_v1.0.csv')[-1]

attach(data)
month_length = c(31,28,31,30,31,30,31,31,30,31,30,31)
head(data)

##   vehicle           model      region season month year distance
## 1 10bc0a12 Nissan Leaf (30 kWh) Coastal Otago Winter     8 2021    2487
## 2 10bc0a12 Nissan Leaf (30 kWh) Coastal Otago Spring     9 2021    1688
## 3 060ffa8d  Hyundai Ioniq (EV)    Waikato Winter     8 2018     880
## 4 060ffa8d  Hyundai Ioniq (EV)    Waikato Spring     9 2018    1352
## 5 060ffa8d  Hyundai Ioniq (EV)    Waikato Spring    10 2018    1869
## 6 060ffa8d  Hyundai Ioniq (EV)    Waikato Spring    11 2018    1286
##   efficiency    kwh
## 1         6.0 414.5
## 2         5.8 291.0
## 3         7.9 111.4
## 4         8.0 169.0
## 5         7.7 242.7
## 6         7.9 162.8

str(data)

## 'data.frame': 23129 obs. of  9 variables:
## $ vehicle    : chr "10bc0a12" "10bc0a12" "060ffa8d" "060ffa8d" ...
## $ model      : chr "Nissan Leaf (30 kWh)" "Nissan Leaf (30 kWh)" "Hyundai Ioniq (EV)" "Hyundai Ioniq (EV)"
## $ region     : chr "Coastal Otago" "Coastal Otago" "Waikato" "Waikato" ...
## $ season     : chr "Winter" "Spring" "Winter" "Spring" ...
## $ month      : int 8 9 8 9 10 11 12 1 2 3 ...
## $ year       : int 2021 2021 2018 2018 2018 2018 2018 2019 2019 2019 ...
## $ distance   : num 2487 1688 880 1352 1869 ...
## $ efficiency: num 6 5.8 7.9 8 7.7 7.9 8 8.26 8.33 8.3 ...
## $ kwh        : num 414 291 111 169 243 ...

summary(data)

##   vehicle           model      region      season
## Length:23129    Length:23129    Length:23129    Length:23129
## Class :character Class :character Class :character Class :character
```

```

##   Mode :character   Mode :character   Mode :character   Mode :character
## 
## 
## 
##       month           year      distance    efficiency
##   Min.   : 1.000   Min.   :2014   Min.   : 0   Min.   : 2.000
##   1st Qu.: 4.000   1st Qu.:2018   1st Qu.: 564   1st Qu.: 6.300
##   Median : 7.000   Median :2019   Median : 919   Median : 6.900
##   Mean   : 6.497   Mean   :2019   Mean   :1077   Mean   : 6.893
##   3rd Qu.: 9.000   3rd Qu.:2020   3rd Qu.:1415   3rd Qu.: 7.500
##   Max.   :12.000   Max.   :2021   Max.   :8695   Max.   :11.100
## 
##       kwh
##   Min.   : 0.0
##   1st Qu.: 80.5
##   Median :132.1
##   Mean   :159.9
##   3rd Qu.:208.3
##   Max.   :3105.4

data %>%
  group_by(year) %>%
  summarise(count = n_distinct(vehicle))

## # A tibble: 8 x 2
##       year count
##   <int> <int>
## 1 2014     1
## 2 2015     5
## 3 2016    23
## 4 2017   287
## 5 2018   755
## 6 2019   850
## 7 2020   728
## 8 2021   637

data %>%
  group_by(region) %>%
  summarise(count = n_distinct(vehicle)) %>%
  arrange(-count)

## # A tibble: 25 x 2
##       region      count
##   <chr>        <int>
## 1 Auckland      334
## 2 Wellington    238
## 3 Christchurch  146
## 4 Coastal Otago 133
## 5 Waikato       65
## 6 Bay of Plenty  53
## 7 North Canterbury 35
## 8 Central Otago  31
## 9 Mid Canterbury 31
## 10 Nelson        31
## # ... with 15 more rows

```

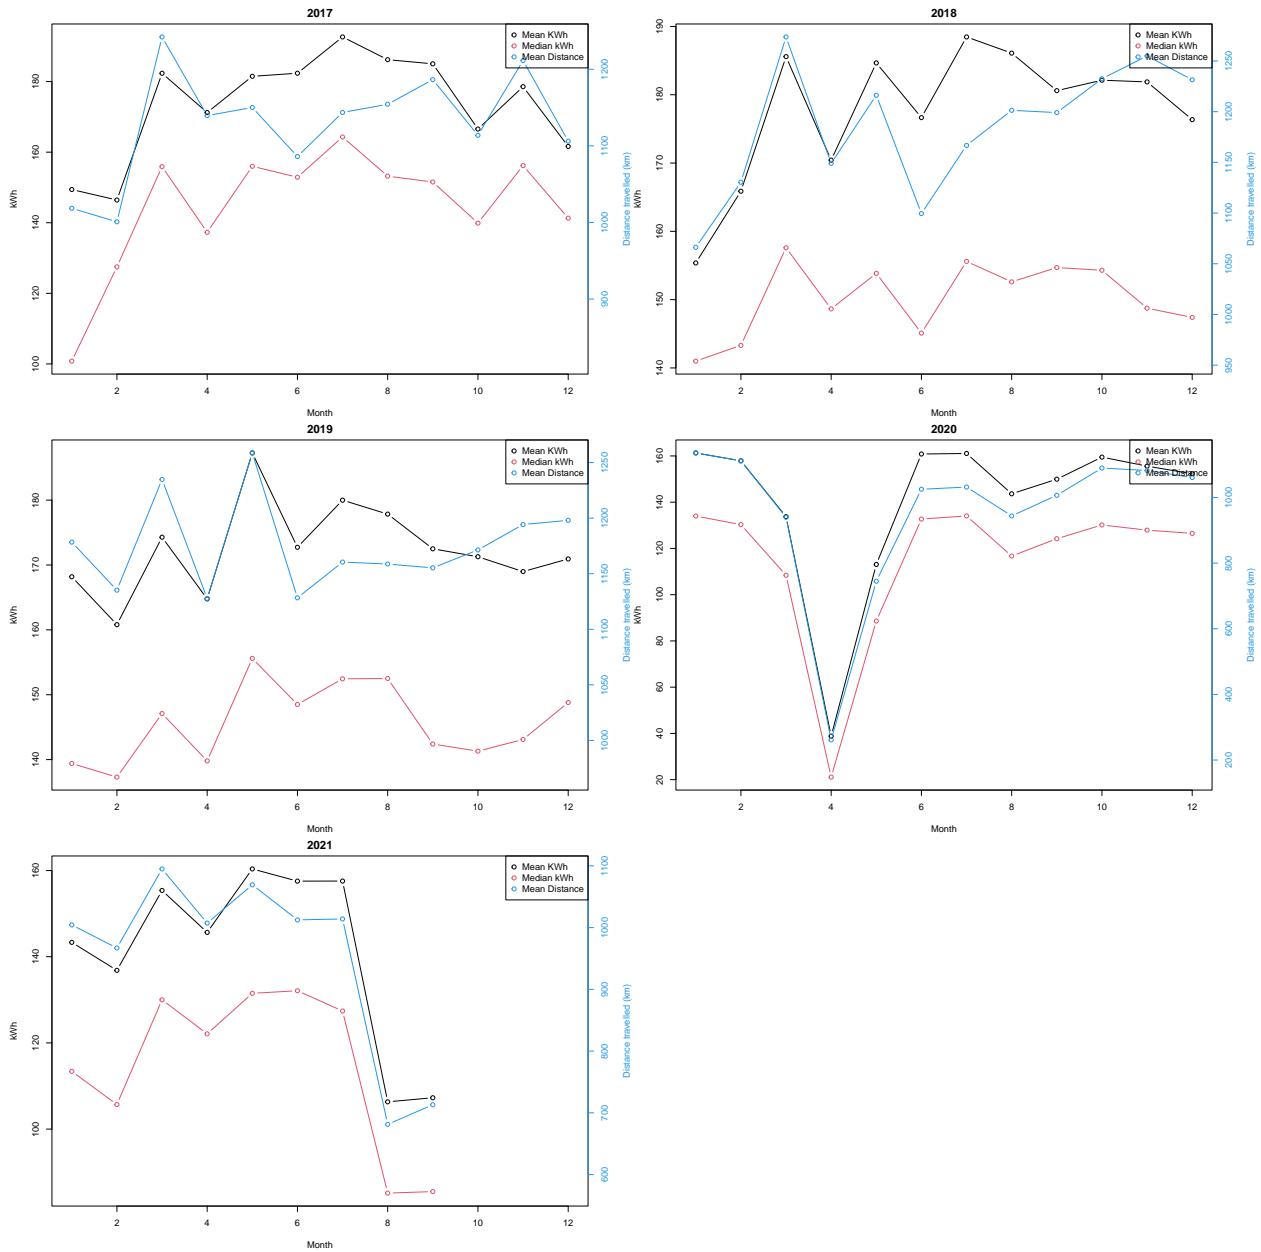
can't really use before 2017

```
mean_data = list()
for (i in 2017:2021) {
  mean_data[[i-2016]] = data[year == i,] %>%
    group_by(month) %>%
    summarise(mean_kwh = mean(kwh), mean_dist = mean(distance), median_kwh = median(kwh), median_dist =
}

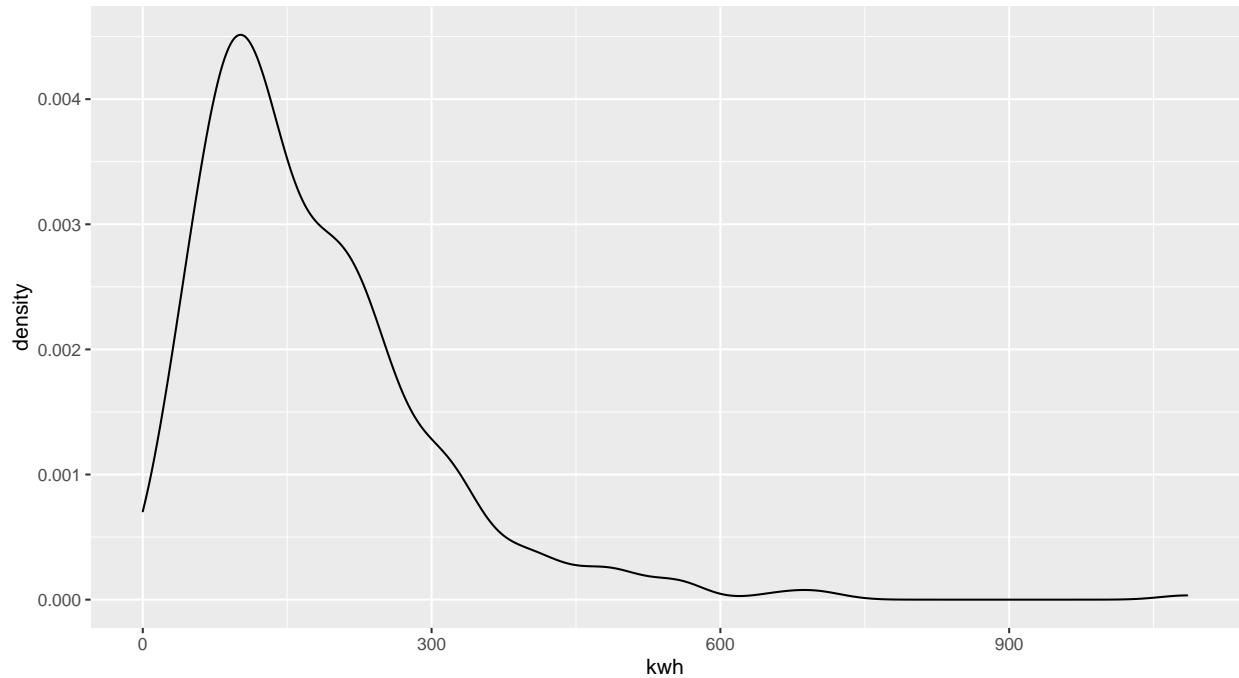
par(mfrow=c(3,2))
for (i in 2017:2021) {
  par(mar = c(4, 4, 2, 4))
  cur_year = data.frame(mean_data[[i-2016]])
  plot(cur_year$month, cur_year$mean_kwh, type = 'b', main = i, xlab = "Month", ylab = "kWh", xlim = c(
    ylim = c(min(c(cur_year$mean_kwh,cur_year$median_kwh)),max(c(cur_year$mean_kwh,cur_year$median_kwh)))
  points(cur_year$month, cur_year$median_kwh, type = 'b', col = 2)

  par(new=TRUE)
  plot(cur_year$month, cur_year$mean_dist, xlab="", ylab="", xlim = c(1,12),
    ylim = c(min(c(cur_year$mean_dist,cur_year$median_dist)),max(c(cur_year$mean_dist,cur_year$median_dist)))
    axes=FALSE, type="b", col=4)
  mtext("Distance travelled (km)",side=4,col=4,line=3, cex = 0.7)
  axis(4, ylim=c(min(c(cur_year$mean_dist,cur_year$median_dist)),max(c(cur_year$mean_dist,cur_year$median_dist)))
    col=4,col.axis=4)
  legend("topright", legend = c("Mean KWh", "Median kWh", "Mean Distance"), col = c(1:2,4), pch = 1)
}

}
```

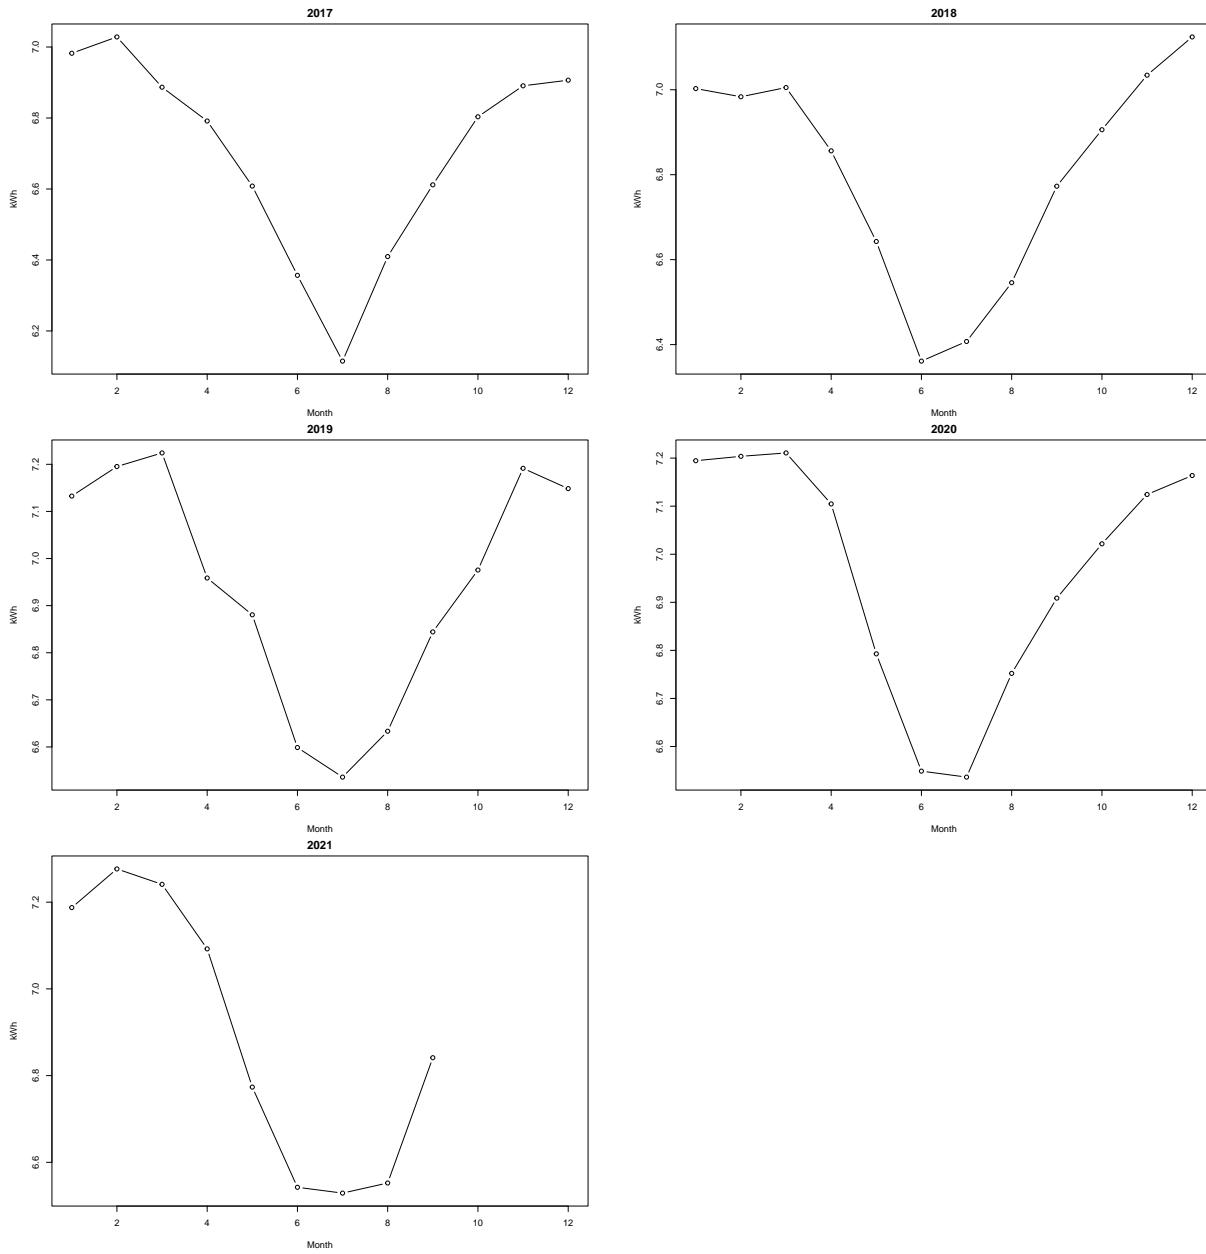


```
ggplot(data = data[year == 2018 & month == 6,], aes(x = kwh)) + geom_density(alpha = 0.5, adjust = 1)
```



```
#ggplot(data = patient, aes(x=age))+ geom_histogram(binwidth =10, color = "peachpuff3", fill = "hotpink")

par(mfrow=c(3,2))
for (i in 2017:2021) {
  par(mar = c(4, 4, 2, 4))
  cur_year = data.frame(mean_data[[i-2016]])
  plot(cur_year$month, cur_year$mean_ef, type = 'b', main = i, xlab = "Month", ylab = "kWh", xlim = c(1,12))
  points(cur_year$month, cur_year$median_kwh, type = 'b', col = 2)
}
```



```

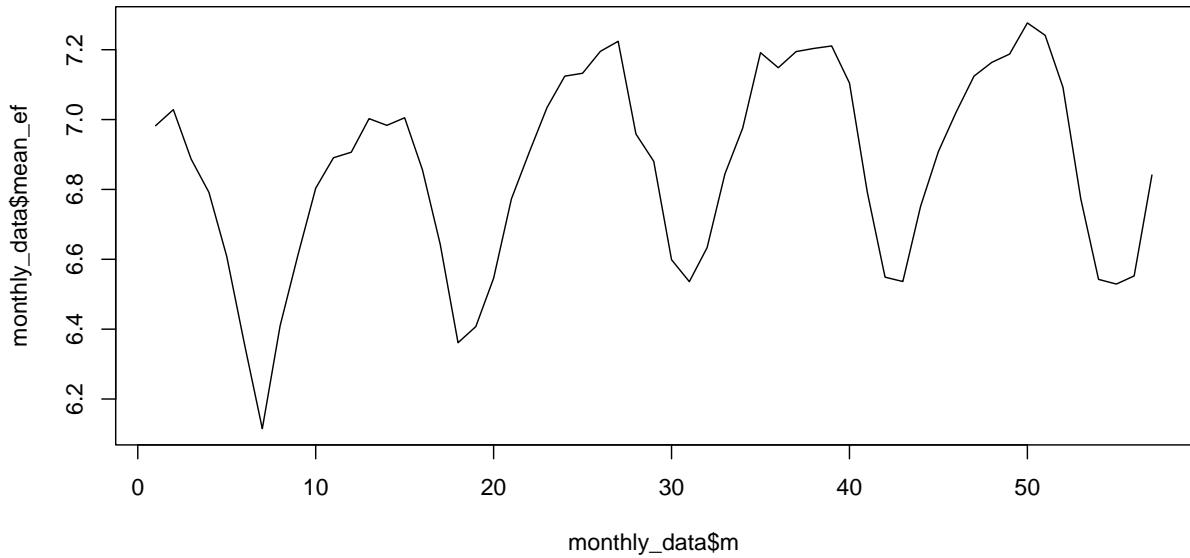
monthly_data = data[year >= 2017, ] %>%
  group_by(year, month) %>%
  summarise(mean_kwh = mean(kwh), mean_dist = mean(distance), median_kwh = median(kwh), median_dist = median(distance))

## `summarise()` has grouped output by 'year'. You can override using the '.groups' argument.

monthly_data$m = 1:nrow(monthly_data)

plot(x = monthly_data$m, y = monthly_data$mean_ef, type = 'l')

```



try a couple of models

```
monthly_eff_lm = lm(mean_ef ~ m+factor(month), data = monthly_data)
summary(monthly_eff_lm)
```

```
##
## Call:
## lm(formula = mean_ef ~ m + factor(month), data = monthly_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -0.169609 -0.035832 -0.002043  0.041524  0.140991 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 6.9540070  0.0361299 192.472 < 2e-16 ***
## m           0.0058376  0.0005981   9.760 1.41e-12 ***
## factor(month)2 0.0316996  0.0465181   0.681 0.499161  
## factor(month)3 0.0019524  0.0465296   0.042 0.966720  
## factor(month)4 -0.1568045  0.0465488  -3.369 0.001579 ** 
## factor(month)5 -0.3837543  0.0465757  -8.239 1.83e-10 ***
## factor(month)6 -0.6475665  0.0466103 -13.893 < 2e-16 ***
## factor(month)7 -0.7102034  0.0466525 -15.223 < 2e-16 ***
## factor(month)8 -0.5621573  0.0467023 -12.037 1.63e-15 ***
## factor(month)9 -0.3508861  0.0467597  -7.504 2.10e-09 *** 
## factor(month)10 -0.1907978  0.0493684  -3.865 0.000362 *** 
## factor(month)11 -0.0631082  0.0493938  -1.278 0.208072  
## factor(month)12 -0.0433899  0.0494264  -0.878 0.384783  
## ---      
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```

## Residual standard error: 0.07355 on 44 degrees of freedom
## Multiple R-squared:  0.9451, Adjusted R-squared:  0.9301
## F-statistic: 63.08 on 12 and 44 DF,  p-value: < 2.2e-16

monthly_eff_qm = lm(mean_ef ~ m+I(m^2)+factor(month),data = monthly_data)
summary(monthly_eff_qm)

##
## Call:
## lm(formula = mean_ef ~ m + I(m^2) + factor(month), data = monthly_data)
##
## Residuals:
##       Min     1Q   Median     3Q    Max 
## -0.136916 -0.035349 -0.000536  0.029554  0.104565
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 6.863e+00 3.339e-02 205.498 < 2e-16 ***
## m           1.571e-02 1.936e-03  8.115 3.26e-10 ***
## I(m^2)      -1.703e-04 3.238e-05 -5.259 4.32e-06 ***
## factor(month)2 3.051e-02 3.671e-02  0.831  0.4105  
## factor(month)3 -9.092e-05 3.672e-02 -0.002  0.9980  
## factor(month)4 -1.594e-01 3.674e-02 -4.338 8.55e-05 ***
## factor(month)5 -3.865e-01 3.676e-02 -10.514 1.84e-13 ***
## factor(month)6 -6.501e-01 3.678e-02 -17.674 < 2e-16 ***
## factor(month)7 -7.122e-01 3.682e-02 -19.346 < 2e-16 ***
## factor(month)8 -5.633e-01 3.686e-02 -15.286 < 2e-16 ***
## factor(month)9 -3.509e-01 3.690e-02 -9.509 3.89e-12 ***
## factor(month)10 -2.117e-01 3.916e-02 -5.407 2.64e-06 ***
## factor(month)11 -8.422e-02 3.918e-02 -2.149  0.0373 *  
## factor(month)12 -6.433e-02 3.921e-02 -1.641  0.1081  
## --- 
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.05804 on 43 degrees of freedom
## Multiple R-squared:  0.9666, Adjusted R-squared:  0.9565
## F-statistic: 95.63 on 13 and 43 DF,  p-value: < 2.2e-16

monthly_eff_lm_log = lm(mean_ef ~ m+I(log(m))+factor(month),data = monthly_data)
summary(monthly_eff_lm_log)

##
## Call:
## lm(formula = mean_ef ~ m + I(log(m)) + factor(month), data = monthly_data)
##
## Residuals:
##       Min     1Q   Median     3Q    Max 
## -0.163993 -0.046093  0.001669  0.037024  0.127306
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 6.880296  0.051541 133.491 < 2e-16 ***
## m           0.003369  0.001392   2.420  0.019840 * 

```

```

## I(log(m))      0.050961  0.026137  1.950 0.057746 .
## factor(month)2 0.025470  0.045217  0.563 0.576164
## factor(month)3 -0.007495  0.045375 -0.165 0.869569
## factor(month)4 -0.168201  0.045511 -3.696 0.000616 ***
## factor(month)5 -0.396378  0.045622 -8.688 5.14e-11 ***
## factor(month)6 -0.660944  0.045712 -14.459 < 2e-16 ***
## factor(month)7 -0.723996  0.045784 -15.813 < 2e-16 ***
## factor(month)8 -0.576107  0.045845 -12.567 5.46e-16 ***
## factor(month)9 -0.364788  0.045896 -7.948 5.61e-10 ***
## factor(month)10 -0.210394  0.048912 -4.302 9.59e-05 ***
## factor(month)11 -0.082660  0.048931 -1.689 0.098399 .
## factor(month)12 -0.062752  0.048942 -1.282 0.206658
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.07131 on 43 degrees of freedom
## Multiple R-squared:  0.9495, Adjusted R-squared:  0.9343
## F-statistic: 62.23 on 13 and 43 DF,  p-value: < 2.2e-16

monthly_eff_qm_log = lm(mean_ef ~ m+I(log(m))+I(m^2)+factor(month), data = monthly_data)
summary(monthly_eff_qm_log)

```

```

##
## Call:
## lm(formula = mean_ef ~ m + I(log(m)) + I(m^2) + factor(month),
##     data = monthly_data)
##
## Residuals:
##       Min        1Q        Median         3Q        Max
## -0.116381 -0.028026 -0.003808  0.031606  0.076696
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 6.977e+00 3.700e-02 188.537 < 2e-16 ***
## m           3.423e-02 4.321e-03   7.921 7.21e-10 ***
## I(log(m))  -1.506e-01 3.267e-02  -4.611 3.73e-05 ***
## I(m^2)      -3.637e-04 4.972e-05  -7.315 5.17e-09 ***
## factor(month)2 4.756e-02 3.049e-02   1.560 0.126276
## factor(month)3 2.551e-02 3.078e-02   0.829 0.411873
## factor(month)4 -1.286e-01 3.101e-02  -4.146 0.000161 ***
## factor(month)5 -3.523e-01 3.120e-02 -11.290 2.65e-14 ***
## factor(month)6 -6.135e-01 3.135e-02 -19.568 < 2e-16 ***
## factor(month)7 -6.738e-01 3.148e-02 -21.405 < 2e-16 ***
## factor(month)8 -5.235e-01 3.159e-02 -16.570 < 2e-16 ***
## factor(month)9 -3.098e-01 3.170e-02  -9.773 2.22e-12 ***
## factor(month)10 -1.776e-01 3.312e-02  -5.362 3.26e-06 ***
## factor(month)11 -5.042e-02 3.313e-02  -1.522 0.135513
## factor(month)12 -3.090e-02 3.313e-02  -0.933 0.356286
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.04785 on 42 degrees of freedom
## Multiple R-squared:  0.9778, Adjusted R-squared:  0.9704
## F-statistic: 132.2 on 14 and 42 DF,  p-value: < 2.2e-16

```

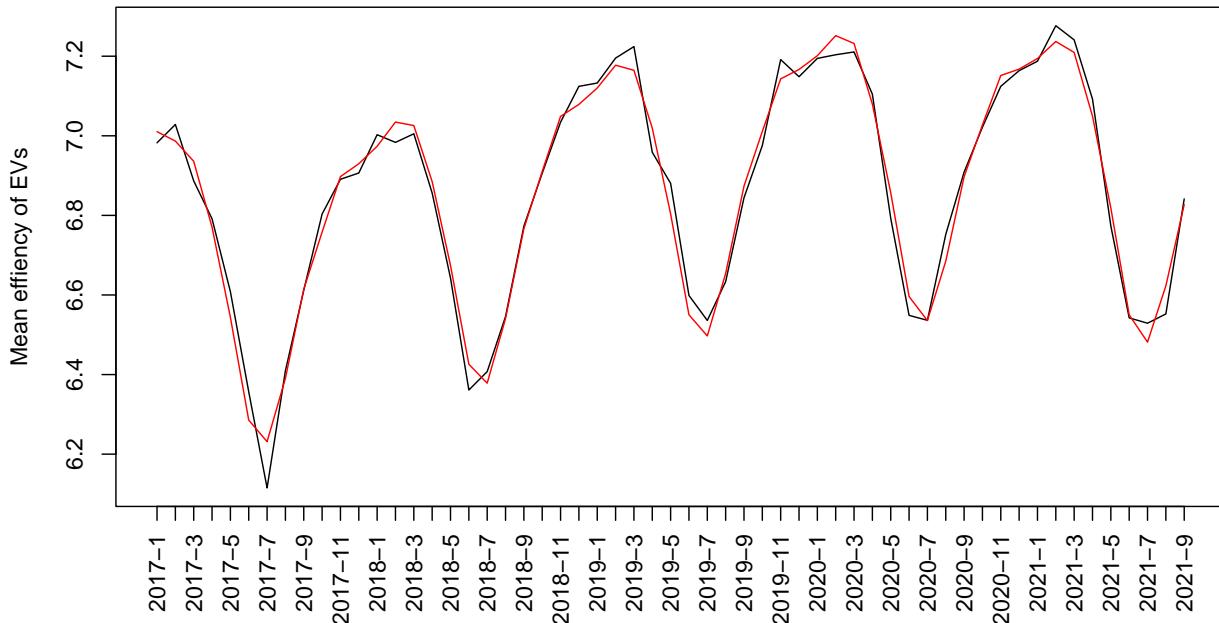
```
anova(monthly_eff_lm, monthly_eff_lm_log, monthly_eff_qm, monthly_eff_qm_log)
```

```
## Analysis of Variance Table
##
## Model 1: mean_ef ~ m + factor(month)
## Model 2: mean_ef ~ m + I(log(m)) + factor(month)
## Model 3: mean_ef ~ m + I(m^2) + factor(month)
## Model 4: mean_ef ~ m + I(log(m)) + I(m^2) + factor(month)
##   Res.Df      RSS Df Sum of Sq    F    Pr(>F)
## 1     44 0.237993
## 2     43 0.218662  1  0.019331  8.4433  0.005824 **
## 3     43 0.144837  0  0.073824
## 4     42 0.096161  1  0.048676 21.2602 3.726e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

par(mar = c(5, 4, 2, 1))
plot(monthly_data$m, monthly_data$mean_ef, type = 'l', xaxt = "n", xlab = "", ylab = "Mean efficiency of EVs", col = "black")
lines(monthly_data$m, predict(monthly_eff_qm_log), col = 'red')

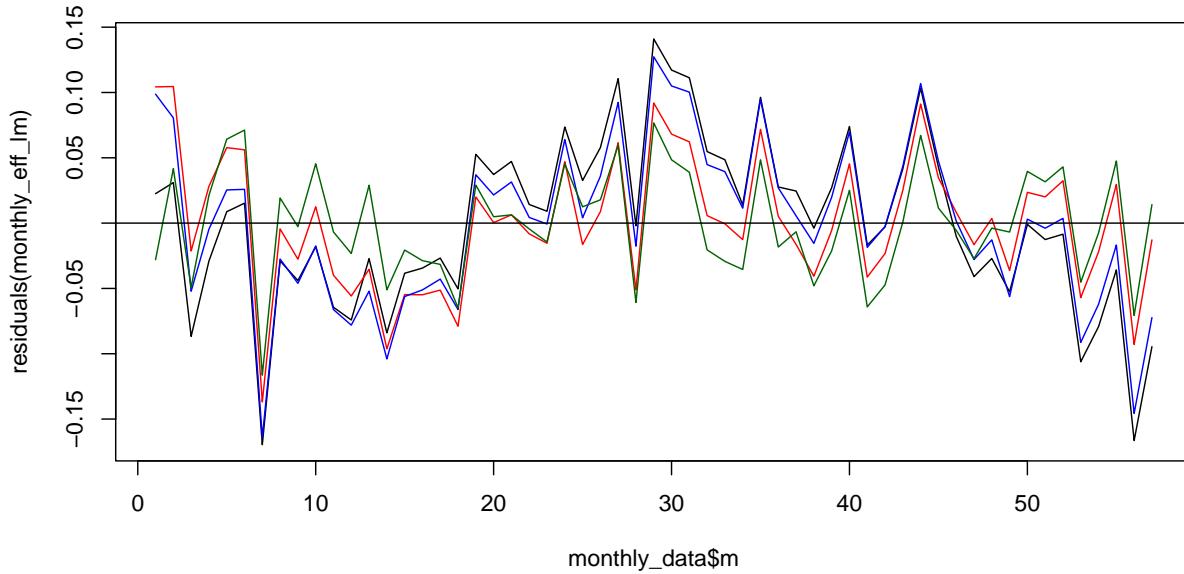
axis(1, labels = paste(monthly_data$year, monthly_data$month, sep = "-"), at = monthly_data$m, las = 2, s
```

Time series of EV efficiencies



```
plot(monthly_data$m, residuals(monthly_eff_lm), type = 'l')
points(monthly_data$m, residuals(monthly_eff_qm), type = 'l', col = "red")
points(monthly_data$m, residuals(monthly_eff_lm_log), type = 'l', col = "blue")
points(monthly_data$m, residuals(monthly_eff_qm_log), type = 'l', col = "darkgreen")
```

```
abline(h= 0)
```



```
eff_series = ts(monthly_data$mean_ef, frequency = 12)
adf.test(eff_series, alternative = "stationary")
```

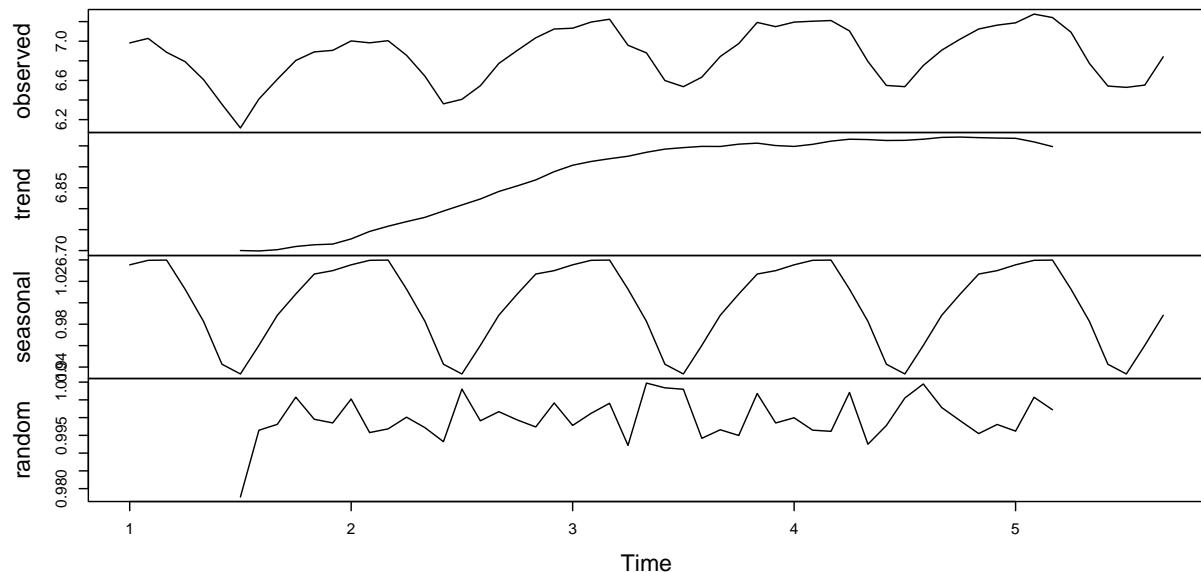
```
## Warning in adf.test(eff_series, alternative = "stationary"): p-value smaller
## than printed p-value

##
## Augmented Dickey-Fuller Test
##
## data: eff_series
## Dickey-Fuller = -4.7051, Lag order = 3, p-value = 0.01
## alternative hypothesis: stationary
```

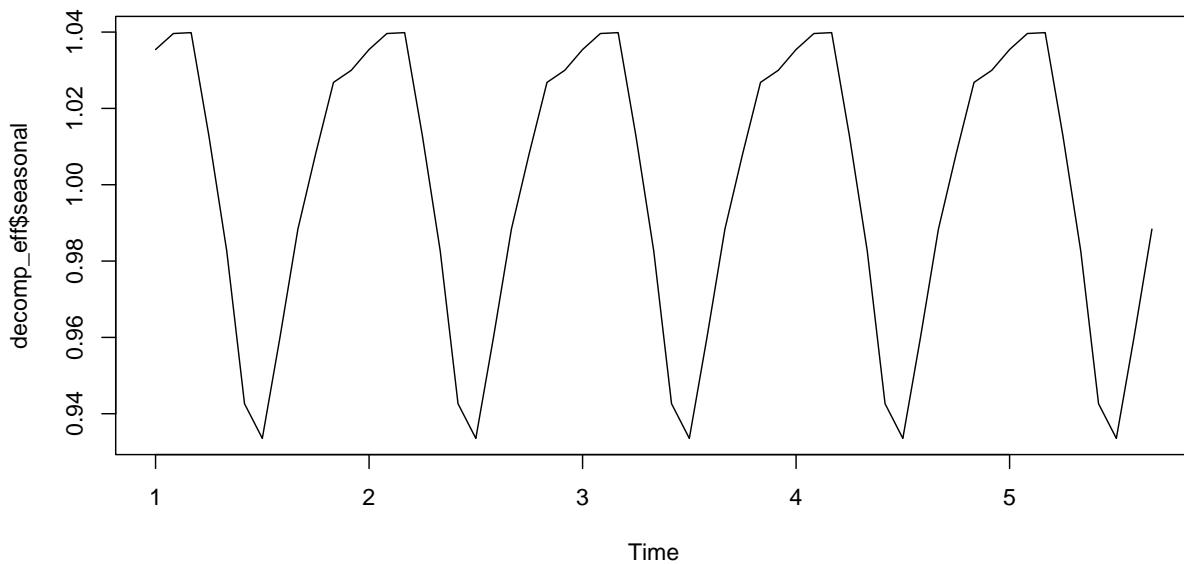
we can reject null hypothesis that data is not-stationary. this makes sense as average efficiency should not have significantly changed in a couple of years. use multiplicative instead of additive as preferable to know estimated extra power use? or should i know total extra power used in season?

```
#decomp_eff = stl(eff_series, s.window = "periodic")
decomp_eff = decompose(eff_series, "multiplicative")
plot(decomp_eff)
```

Decomposition of multiplicative time series



```
plot(decomp_eff$seasonal)
```



```
eff_model = auto.arima(eff_series) #uses AIC to get best model  
summary(eff_model)
```

```
## Series: eff_series  
## ARIMA(1,1,2)(1,1,0)[12]
```

```

##  

## Coefficients:  

##          ar1      ma1      ma2      sar1  

##        0.8306 -1.7118  0.8679 -0.5824  

##  s.e.  0.1549  0.1391  0.1332  0.1348  

##  

## sigma^2 estimated as 0.004102:  log likelihood=56.94  

## AIC=-103.88  AICc=-102.3  BIC=-94.96  

##  

## Training set error measures:  

##               ME      RMSE      MAE      MPE      MAPE      MASE  

## Training set -0.004027656 0.05365432 0.03802298 -0.05547605 0.5542138 0.4152625  

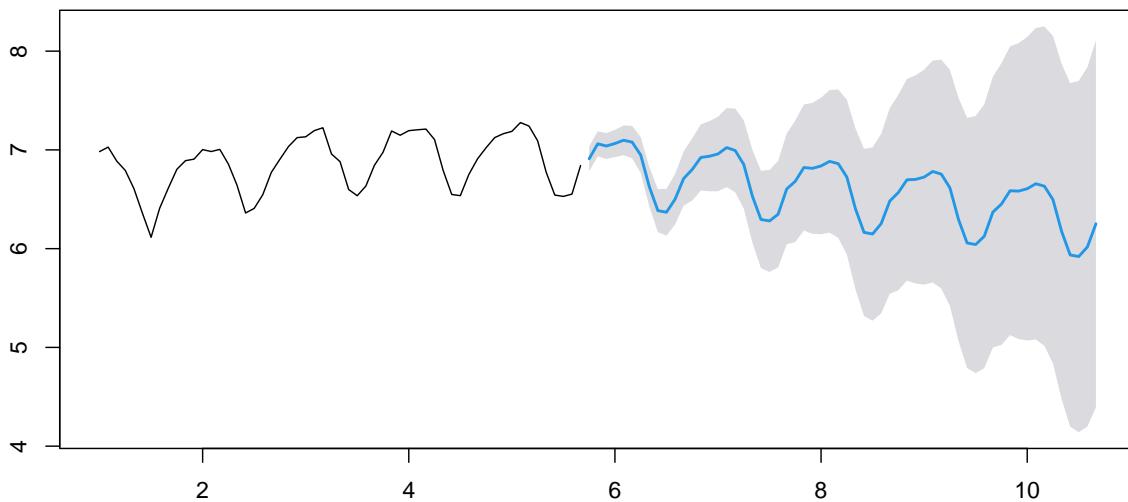
##  

## ACF1  

## Training set -0.02594167
efficiency_forcast = forecast(eff_model, level = 95, h = 60)
plot(efficiency_forcast)

```

Forecasts from ARIMA(1,1,2)(1,1,0)[12]



```
Box.test(efficiency_forcast$resid, lag=6, type="Ljung-Box")
```

```

##  

## Box-Ljung test  

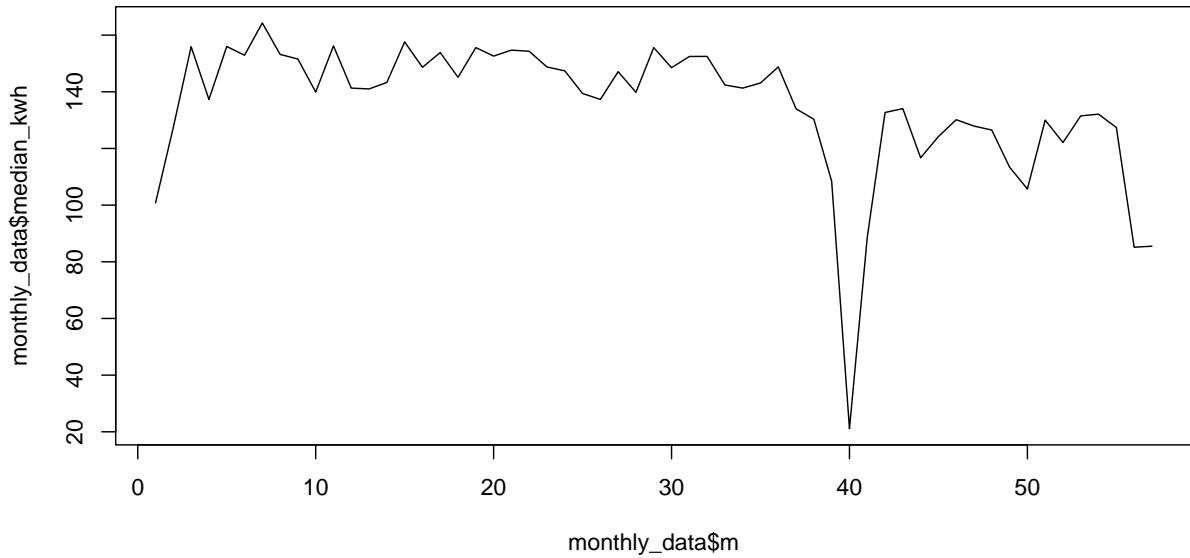
##  

## data: efficiency_forcast$resid  

## X-squared = 2.9725, df = 6, p-value = 0.8123

```

```
plot(x = monthly_data$m, y = monthly_data$median_kwh, type = 'l')
```



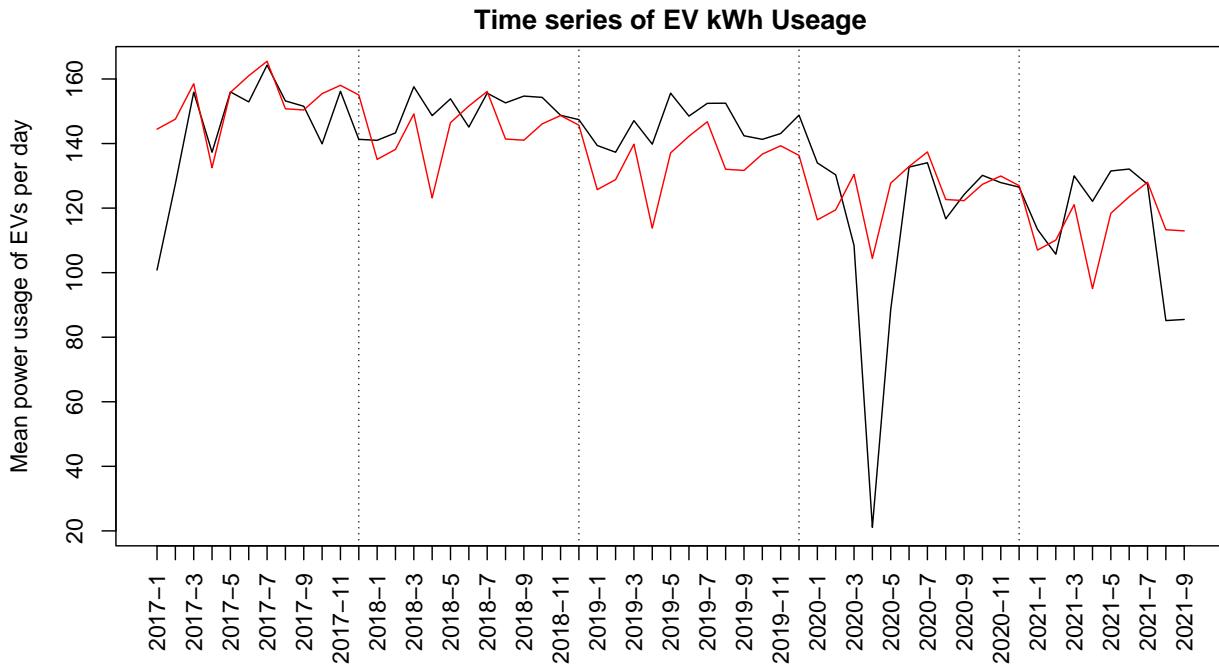
```
monthly_kwh_lm = lm(median_kwh ~ m+factor(month), data = monthly_data)
summary(monthly_kwh_lm)
```

```
##
## Call:
## lm(formula = median_kwh ~ m + factor(month), data = monthly_data)
##
## Residuals:
##    Min      1Q  Median      3Q     Max 
## -83.324 -2.632   2.787   8.572  27.042 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 145.2329   10.0324 14.476 < 2e-16 ***
## m          -0.7805    0.1661 -4.700 2.58e-05 ***
## factor(month)2  3.8805  12.9169  0.300  0.7653  
## factor(month)3 15.6410  12.9201  1.211  0.2325  
## factor(month)4 -9.5885  12.9255 -0.742  0.4621  
## factor(month)5 14.5121  12.9329  1.122  0.2679  
## factor(month)6 20.4426  12.9425  1.579  0.1214  
## factor(month)7 25.7231  12.9543  1.986  0.0533 .  
## factor(month)8 11.7736  12.9681  0.908  0.3689  
## factor(month)9 12.1941  12.9840  0.939  0.3528  
## factor(month)10 18.0340  13.7084  1.316  0.1951  
## factor(month)11 21.3896  13.7154  1.560  0.1260  
## factor(month)12 19.1826  13.7245  1.398  0.1692  
## ---      
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 20.42 on 44 degrees of freedom
## Multiple R-squared:  0.4311, Adjusted R-squared:  0.276
```

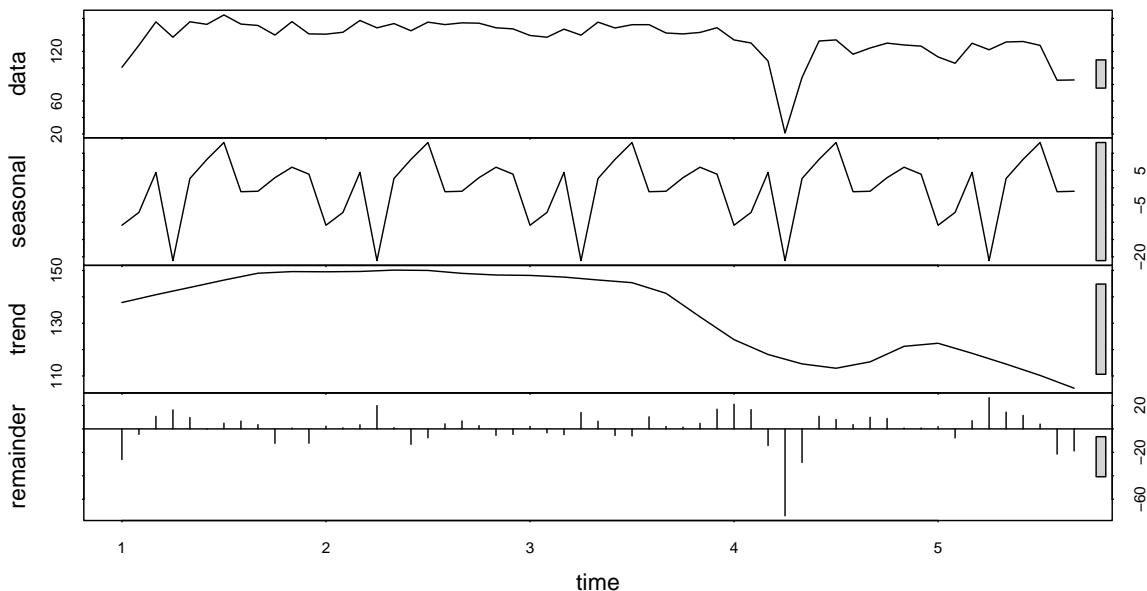
```
## F-statistic: 2.779 on 12 and 44 DF, p-value: 0.006676
```

```
par(mar = c(5, 4, 2, 1))
plot(monthly_data$m, monthly_data$median_kwh, type = 'l', xaxt = "n", xlab = "", ylab = "Mean power usage per day")
lines(monthly_data$m, predict(monthly_kwh_lm), col = 'red')

axis(1, labels = paste(monthly_data$year, monthly_data$month, sep = "-"), at = monthly_data$m, las = 2)
abline(v = 12, lty = 3)
abline(v = 24, lty = 3)
abline(v = 36, lty = 3)
abline(v = 48, lty = 3)
abline(v = 60, lty = 3)
```



```
kwh_series = ts(monthly_data$median_kwh, frequency = 12)
decomp_kwh = stl(kwh_series, s.window = "periodic")
plot(decomp_kwh)
```



```
adf.test(kwh_series, alternative = "stationary")
```

```
##
##  Augmented Dickey-Fuller Test
##
## data: kwh_series
## Dickey-Fuller = -3.5116, Lag order = 3, p-value = 0.04852
## alternative hypothesis: stationary
```

linear model time series but grouped by car model. turns out is basically useless and get more confident on the monthly effiencys by just using the average

```
#monthly_data_model = data[year >= 2017,] %>%
#  group_by(year, month, model) %>%
#  summarise(mean_kwh = mean(kwh), mean_dist = mean(distance), median_kwh = median(kwh), median_dist =
#monthly_data_model$m = as.numeric(factor(paste(monthly_data_model$year,monthly_data_model$month, sep =
#                                         levels = unique(paste(monthly_data_model$year,monthly_data_model$month))))
```



```
#monthly_eff_model_lm = lm(mean_ef ~ m+factor(month)+factor(model),data = monthly_data_model)
#summary(monthly_eff_model_lm)
```

```
fuel_data = read.csv('Dunedin Fuel research/Dunedin fuel usage.csv')
head(fuel_data)
```

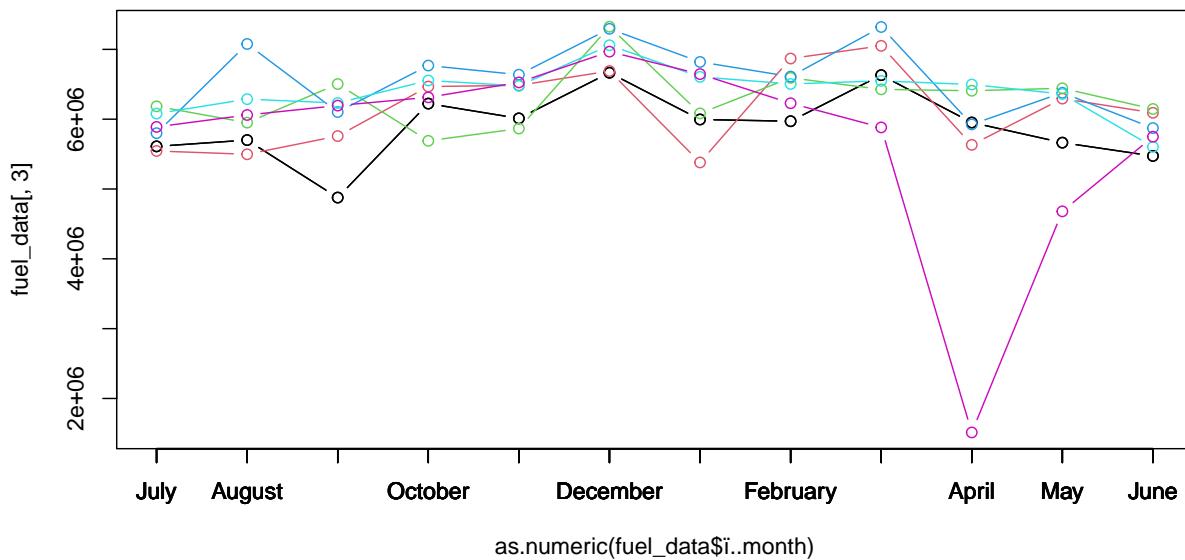
```
##    i..month month_no   X2015   X2016   X2017   X2018   X2019   X2020
## 1      July       7 5610642 5544497 6184078 5799777 6079678 5889484
## 2     August      8 5700740 5497344 5950339 7077168 6286820 6059895
## 3 September     9 4877210 5757994 6503403 6102254 6231719 6194753
## 4 October      10 6223226 6467628 5690018 6768411 6554627 6312884
## 5 November     11 6011595 6484301 5865678 6637329 6476736 6526996
## 6 December     12 6665164 6689692 7324506 7294138 7058879 6967167
```

```

fuel_data$i..month = factor(fuel_data$i..month, levels = fuel_data$i..month)

plot(as.numeric(fuel_data$i..month), fuel_data[,3], type = 'b', col = 1 ,xaxt = "n",
      ylim = c(min(fuel_data[,3:8]),max(fuel_data[,3:8])))
for (i in 3:8) {
  points(as.numeric(fuel_data$i..month), fuel_data[,i], type = 'b', col = i-2 ,xaxt = "n")
  axis(1, labels = as.character(fuel_data$i..month), at = as.numeric(fuel_data$i..month))
}

```



```

mean_fuel = fuel_data %>%
  gather(key = "year", value = "Fuel_used", c("X2017","X2018","X2019")) %>%
  group_by(i..month) %>%
  summarise(average_fuel_use = mean(Fuel_used))

power = c()
for (i in 1:12) {
  #average power useage (kWh/km)
  power[i] = mean(1/data[month==i & year >= 2017 & year <= 2019,]$efficiency, na.rm = T)
}
mean_fuel$power_usage = power[c(7:12,1:6)]
mean_fuel

```

```

## # A tibble: 12 x 3
##   i..month  average_fuel_use power_usage
##   <fct>          <dbl>       <dbl>
## 1 July        6021178.     0.159
## 2 August      6438109.     0.155
## 3 September   6279125.     0.150
## 4 October     6337685.     0.147

```

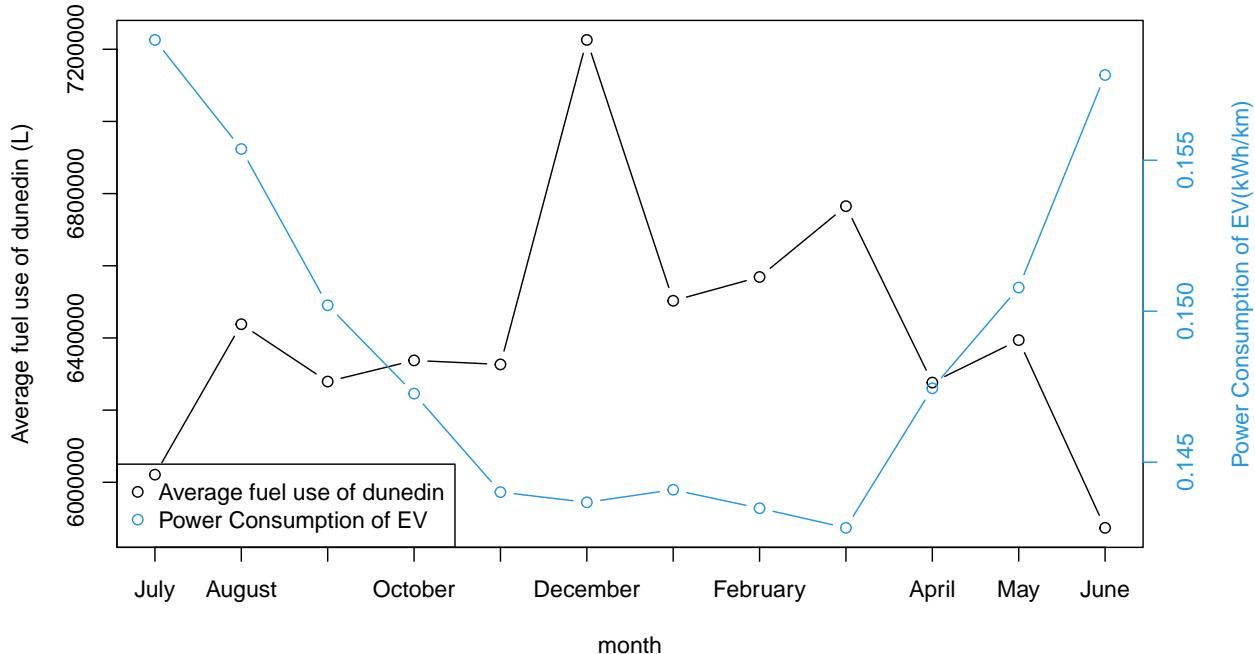
```

## 5 November      6326581.      0.144
## 6 December      7225841.      0.144
## 7 January       6503050.      0.144
## 8 February      6568882.      0.143
## 9 March         6765136.      0.143
## 10 April        6276348.      0.147
## 11 May          6394003.      0.151
## 12 June         5873758.      0.158

par(mar = c(4, 4, 2, 4))
plot(as.numeric(mean_fuel$i..month), mean_fuel$average_fuel_use, type = 'b', xaxt = "n",
     ylab = "Average fuel use of dunedin (L)", xlab = "month")
axis(1, labels = as.character(fuel_data$i..month), at = as.numeric(fuel_data$i..month))

par(new=TRUE)
plot(as.numeric(mean_fuel$i..month), mean_fuel$power_usage, xlab="", ylab="", xlim = c(1,12),
     ylim=c(min(mean_fuel$power_usage),max(mean_fuel$power_usage)),
     axes=FALSE, type="b", col=4)
mtext("Power Consumption of EV(kWh/km)", side=4, col=4, line=3, cex = 1)
axis(4, ylim=c(min(mean_fuel$power_usage),max(mean_fuel$power_usage)),
     col=4, col.axis=4)
legend("bottomleft", legend = c("Average fuel use of dunedin", "Power Consumption of EV"), col = c(1, 4))

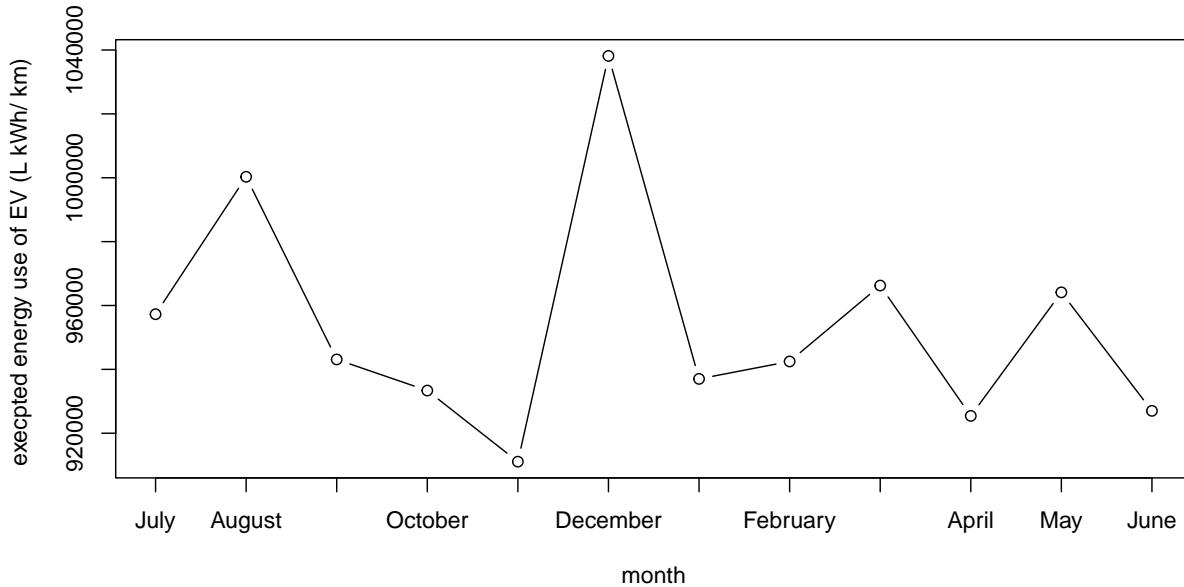
```



```

plot(as.numeric(mean_fuel$i..month), mean_fuel$average_fuel_use * mean_fuel$power_usage, type = 'b', xaxt = "n",
     ylab = "execped energy use of EV (L kWh/ km) ", xlab = "month")
axis(1, labels = as.character(fuel_data$i..month), at = as.numeric(fuel_data$i..month))

```



numbers could probably put to some meaningful scale if use some scaler (L/km) like fuel efficiency of gas car then can just show kwh usage

```
temp_data = read.csv('weather_data/auckland_motat_ews_data.csv', stringsAsFactors = T)
temp_data$Month = factor(temp_data$Month, levels = unique(temp_data$Month))
head(temp_data, 10)
```

```
##      City     Station Year Month Day Hour Temp Twet.C. RH... Tdew.C.
## 1 Auckland Motat Ews 2017   Jan    1    0 15.2   13.7  84.0   12.5
## 2 Auckland Motat Ews 2017   Jan    1    1 15.1   13.8  86.0   12.8
## 3 Auckland Motat Ews 2017   Jan    1    2 14.3   13.3  89.0   12.5
## 4 Auckland Motat Ews 2017   Jan    1    3 14.2   13.4  91.0   12.8
## 5 Auckland Motat Ews 2017   Jan    1    4 13.2   12.5  92.0   11.9
## 6 Auckland Motat Ews 2017   Jan    1    5 12.3   11.8  94.0   11.4
## 7 Auckland Motat Ews 2017   Jan    1    6 13.8   13.3  94.0   12.9
## 8 Auckland Motat Ews 2017   Jan    1    7 15.5   14.9  93.0   14.4
## 9 Auckland Motat Ews 2017   Jan    1    8 18.7   16.7  81.0   15.4
## 10 Auckland Motat Ews 2017   Jan   1    9 21.0   17.1  66.0   14.4
```

```
HDD = function (data, temp = 18) {
  mon_sum = data.frame(Year = NA, Month = NA, HDD = NA)[-1,]
  for (y in min(data$Year):max(data$Year)){
    for (m in unique(data[data$Year == y,]$Month)) {
      #print(paste(y, " ", m))
      adj_temps = -data[data$Year == y & data$Month == m,]$Temp+temp
      #print(sum(adj_temps*(adj_temps > 0))/24)
      mon_sum[nrow(mon_sum) + 1,] = c(y,m, sum(adj_temps*(adj_temps > 0))/24)
    }
  }
  mon_sum$Year = as.factor(mon_sum$Year)
  mon_sum$Month = factor(mon_sum$Month, levels = month.abb)
```

```

mon_sum$HDD = as.numeric(mon_sum$HDD)
return(mon_sum)
}

CDD = function (data, temp = 16) {
  mon_sum = data.frame(Year = NA, Month = NA, CDD = NA)[-1,]
  for (y in min(data$Year):max(data$Year)){
    for (m in levels(data$Month)) {
      #print(paste(y, " ", m))
      adj_temps = data[data$Year == y & data$Month == m,]$Temp-temp
      #print(sum(adj_temps*(adj_temps > 0))/24)
      mon_sum[nrow(mon_sum) + 1,] = c(y,m, sum(adj_temps*(adj_temps > 0))/24)

    }
  }
  mon_sum$Year = as.factor(mon_sum$Year)
  mon_sum$Month = as.factor(mon_sum$Month)
  mon_sum$CDD = as.numeric(mon_sum$CDD)
  return(mon_sum)
}

```

```
HDD_auckland = HDD(temp_data)
```

```
HDD_auckland
```

	##	Year	Month	HDD
## 1	2017	Jan	33.479167	
## 2	2017	Feb	18.050000	
## 3	2017	Mar	23.720833	
## 4	2017	Apr	59.337500	
## 5	2017	May	151.545833	
## 6	2017	Jun	186.275000	
## 7	2017	Jul	231.037500	
## 8	2017	Aug	177.429167	
## 9	2017	Sep	149.337500	
## 10	2017	Oct	101.175000	
## 11	2017	Nov	62.670833	
## 12	2017	Dec	15.516667	
## 13	2018	Jan	2.141667	
## 14	2018	Feb	13.783333	
## 15	2018	Mar	25.395833	
## 16	2018	Apr	61.908333	
## 17	2018	May	131.304167	
## 18	2018	Jun	203.704167	
## 19	2018	Jul	226.125000	
## 20	2018	Aug	201.520833	
## 21	2018	Sep	157.737500	
## 22	2018	Oct	133.100000	
## 23	2018	Nov	76.700000	
## 24	2018	Dec	21.508333	
## 25	2019	Jan	6.995833	
## 26	2019	Feb	15.508333	
## 27	2019	Mar	24.925000	

```

## 28 2019 Apr 97.870833
## 29 2019 May 113.637500
## 30 2019 Jun 236.512500
## 31 2019 Jul 235.900000
## 32 2019 Aug 231.629167
## 33 2019 Sep 167.445833
## 34 2019 Oct 119.800000
## 35 2019 Nov 60.891667
## 36 2019 Dec 30.129167
## 37 2020 Jan 26.433333
## 38 2020 Feb 15.587500
## 39 2020 Mar 48.300000
## 40 2020 Apr 83.641667
## 41 2020 May 136.325000
## 42 2020 Jun 155.308333
## 43 2020 Jul 204.429167
## 44 2020 Aug 188.341667
## 45 2020 Sep 157.812500
## 46 2020 Oct 99.979167
## 47 2020 Nov 58.900000
## 48 2020 Dec 35.283333
## 49 2021 Jan 20.745833
## 50 2021 Feb 18.687500
## 51 2021 Mar 29.212500
## 52 2021 Apr 70.658333
## 53 2021 May 122.470833
## 54 2021 Jun 141.325000
## 55 2021 Jul 215.991667
## 56 2021 Aug 186.450000
## 57 2021 Sep 155.995833
## 58 2021 Oct 91.454167
## 59 2021 Nov 26.754167

```

```
CDD(temp_data, 20)
```

```

##   Year Month      CDD
## 1 2017 Jan 16.85000000
## 2 2017 Feb 33.18750000
## 3 2017 Mar 19.15833333
## 4 2017 Apr 8.70000000
## 5 2017 May 0.11666667
## 6 2017 Jun 0.00000000
## 7 2017 Jul 0.00000000
## 8 2017 Aug 0.00000000
## 9 2017 Sep 0.01666667
## 10 2017 Oct 0.20000000
## 11 2017 Nov 6.89583333
## 12 2017 Dec 36.88750000
## 13 2018 Jan 65.74166667
## 14 2018 Feb 44.68333333
## 15 2018 Mar 33.15833333
## 16 2018 Apr 7.00000000
## 17 2018 May 1.36666667
## 18 2018 Jun 0.00000000

```

```

## 19 2018 Jul 0.0000000
## 20 2018 Aug 0.0000000
## 21 2018 Sep 0.0250000
## 22 2018 Oct 0.37083333
## 23 2018 Nov 3.75833333
## 24 2018 Dec 14.85416667
## 25 2019 Jan 49.92916667
## 26 2019 Feb 46.15833333
## 27 2019 Mar 36.49583333
## 28 2019 Apr 3.09583333
## 29 2019 May 0.58333333
## 30 2019 Jun 0.00000000
## 31 2019 Jul 0.00000000
## 32 2019 Aug 0.00000000
## 33 2019 Sep 0.00000000
## 34 2019 Oct 0.29166667
## 35 2019 Nov 10.22083333
## 36 2019 Dec 20.06250000
## 37 2020 Jan 33.60833333
## 38 2020 Feb 54.75416667
## 39 2020 Mar 21.21250000
## 40 2020 Apr 5.32083333
## 41 2020 May 0.14166667
## 42 2020 Jun 0.00000000
## 43 2020 Jul 0.00000000
## 44 2020 Aug 0.00000000
## 45 2020 Sep 0.00000000
## 46 2020 Oct 2.52500000
## 47 2020 Nov 4.62500000
## 48 2020 Dec 15.27500000
## 49 2021 Jan 33.09166667
## 50 2021 Feb 33.62083333
## 51 2021 Mar 18.38750000
## 52 2021 Apr 5.70416667
## 53 2021 May 0.15416667
## 54 2021 Jun 0.00000000
## 55 2021 Jul 0.00000000
## 56 2021 Aug 0.00000000
## 57 2021 Sep 0.06666667
## 58 2021 Oct 0.26666667
## 59 2021 Nov 6.37083333
## 60 2021 Dec 0.00000000

mean_HDD_auckland = HDD_auckland %>%
  group_by(Month) %>%
  summarise(average_HDD = mean(HDD))

power = c()
eff = c()
kwh = c()
km = c()
for (i in 1:12) {
  #average power usage (kWh/km)
  power[i] = mean(1/data[month==i & region == "Auckland" & year >= 2017,]$efficiency, na.rm = T)
}

```

```

eff[i] = mean(data[month==i & region == "Auckland" & year >= 2017,]$efficiency, na.rm = T)
kwh[i] = mean(data[month==i & region == "Auckland" & year >= 2017,]$kwh, na.rm = T)
km[i] = mean(data[month==i & region == "Auckland" & year >= 2017,]$distance, na.rm = T)

}

mean_HDD_auckland$power_usage = power
mean_HDD_auckland$efficiency = eff
mean_HDD_auckland$kwh = kwh
mean_HDD_auckland$km = km
mean_HDD_auckland

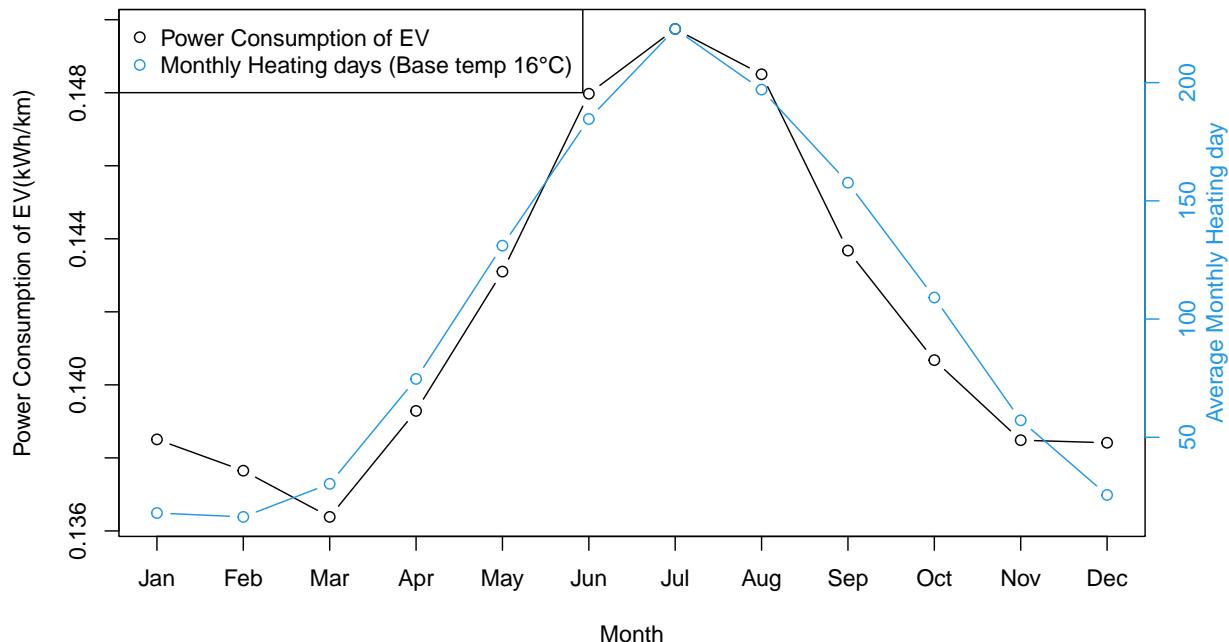
## # A tibble: 12 x 6
##   Month average_HDD power_usage efficiency     kwh      km
##   <fct>     <dbl>      <dbl>       <dbl>    <dbl>    <dbl>
## 1 Jan        18.0      0.139      7.32    156.  1126.
## 2 Feb        16.3      0.138      7.38    149.  1074.
## 3 Mar        30.3      0.136      7.44    155.  1134.
## 4 Apr        74.7      0.139      7.32    123.   883.
## 5 May       131.       0.143      7.14    154.  1078.
## 6 Jun       185.       0.148      6.90    162.  1089.
## 7 Jul       223.       0.150      6.84    166.  1124.
## 8 Aug       197.       0.149      6.89    146.   983.
## 9 Sep       158.       0.144      7.09    138.   960.
## 10 Oct      109.       0.141      7.22    160.  1137.
## 11 Nov      57.2       0.138      7.34    167.  1202.
## 12 Dec      25.6       0.138      7.33    166.  1193.

par(mar = c(4, 4, 2, 4))
plot(as.numeric(mean_HDD_auckland$Month), mean_HDD_auckland$power_usage, type = 'b', xaxt = "n",
     ylab = "Power Consumption of EV(kWh/km)", xlab = "Month", main = "Auckland EV power consumption and monthly heating days", 
     axis(1, labels = as.character(mean_HDD_auckland$Month)), at = as.numeric(mean_HDD_auckland$Month))

par(new=TRUE)
plot(as.numeric(mean_HDD_auckland$Month), mean_HDD_auckland$average_HDD, xlab="", ylab="", xlim = c(1,12),
     ylim=c(min(mean_HDD_auckland$average_HDD),max(mean_HDD_auckland$average_HDD)),
     axes=FALSE, type="b", col=4)
mtext("Average Monthly Heating day", side=4, col=4, line=2, cex = 1)
axis(4, ylim=c(min(mean_HDD_auckland$average_HDD),max(mean_HDD_auckland$average_HDD)),
     col=4, col.axis=4)
legend("topleft", legend = c( "Power Consumption of EV", "Monthly Heating days (Base temp 16°C)"), col = 4)

```

Auckland EV power consumption and Monthly Heating Days by Month

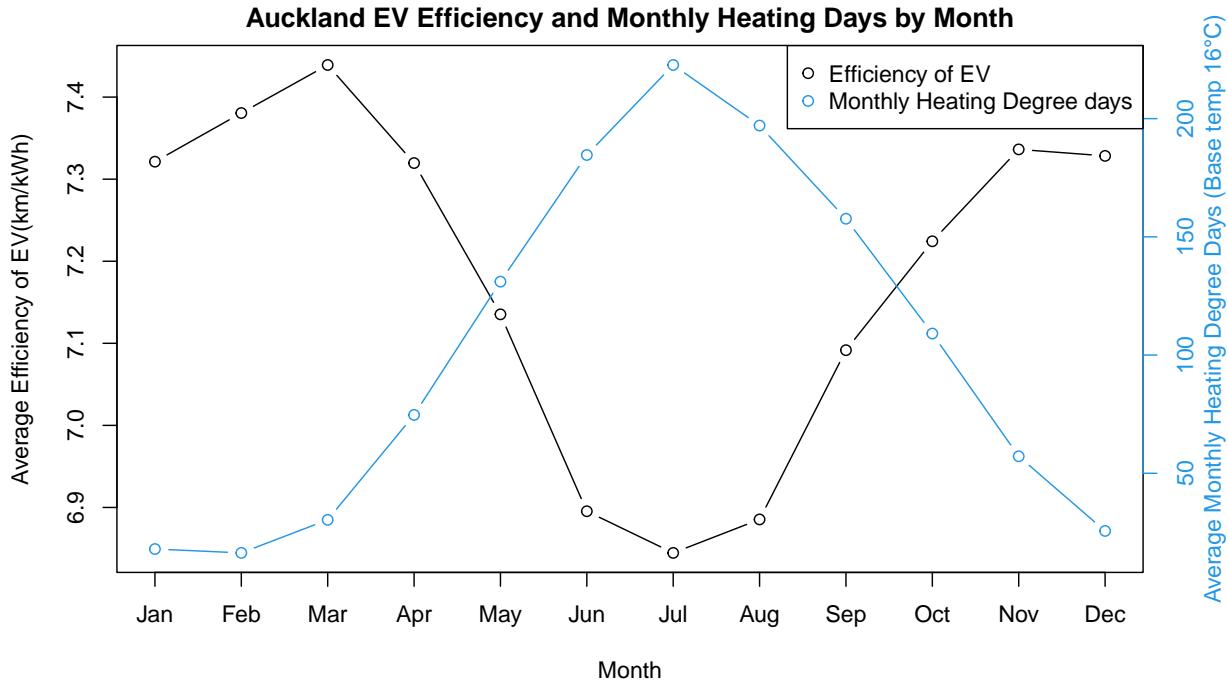


```

par(mar = c(4, 4, 2, 4))
plot(as.numeric(mean_HDD_auckland$Month), mean_HDD_auckland$efficiency, type = 'b', xaxt = "n",
     ylab = "Average Efficiency of EV(km/kWh)", xlab = "Month", main = "Auckland EV Efficiency and Mont",
axis(1, labels = as.character(mean_HDD_auckland$Month), at = as.numeric(mean_HDD_auckland$Month))

par(new=TRUE)
plot(as.numeric(mean_HDD_auckland$Month), mean_HDD_auckland$average_HDD, xlab="", ylab="", xlim = c(1,12),
     ylim=c(min(mean_HDD_auckland$average_HDD),max(mean_HDD_auckland$average_HDD)),
     axes=FALSE, type="b", col=4)
mtext("Average Monthly Heating Degree Days (Base temp 16°C)", side=4, col=4, line=2, cex = 1)
axis(4, ylim=c(min(mean_HDD_auckland$average_HDD),max(mean_HDD_auckland$average_HDD)),
     col=4, col.axis=4)
legend("topright", legend = c( "Efficiency of EV", "Monthly Heating Degree days"), col = c(1, 4), pch =

```



linear model of efficiency and HDD with averages from 2017 to 2021 in Auckland

```
avg_eff_HDD_auckland = data[data$year >= 2017 & region == "Auckland",] %>%
  group_by(year, month) %>%
  summarise(average_efficiency = mean(efficiency), avg_kwh = mean(kwh))
```

'summarise()' has grouped output by 'year'. You can override using the '.groups' argument.

```
#avg_eff_HDD_auckland$HDD = HDD_auckland$HDD
HDD_auckland$month = as.numeric(HDD_auckland$Month)
avg_eff_HDD_auckland = merge(avg_eff_HDD_auckland, HDD_auckland, by.x = c("year", "month"), by.y = c("Year", "Month"))
head(avg_eff_HDD_auckland, 10)
```

	year	month	efficiency	avg_kwh	Month	HDD
## 1	2017	1	6.890000	127.7750	Jan	33.47917
## 2	2017	2	7.066250	124.9750	Feb	18.05000
## 3	2017	3	6.798182	181.1636	Mar	23.72083
## 4	2017	4	7.286154	193.3000	Apr	59.33750
## 5	2017	5	7.053333	189.0933	May	151.54583
## 6	2017	6	6.932941	175.2000	Jun	186.27500
## 7	2017	7	6.794167	185.2792	Jul	231.03750
## 8	2017	8	6.835128	161.5923	Aug	177.42917
## 9	2017	9	6.985500	159.9525	Sep	149.33750
## 10	2017	10	7.069434	163.4340	Oct	101.17500

```
all_lm = lm(data = avg_eff_HDD_auckland, efficiency ~ HDD)
summary(all_lm)
```

```
##
## Call:
```

```

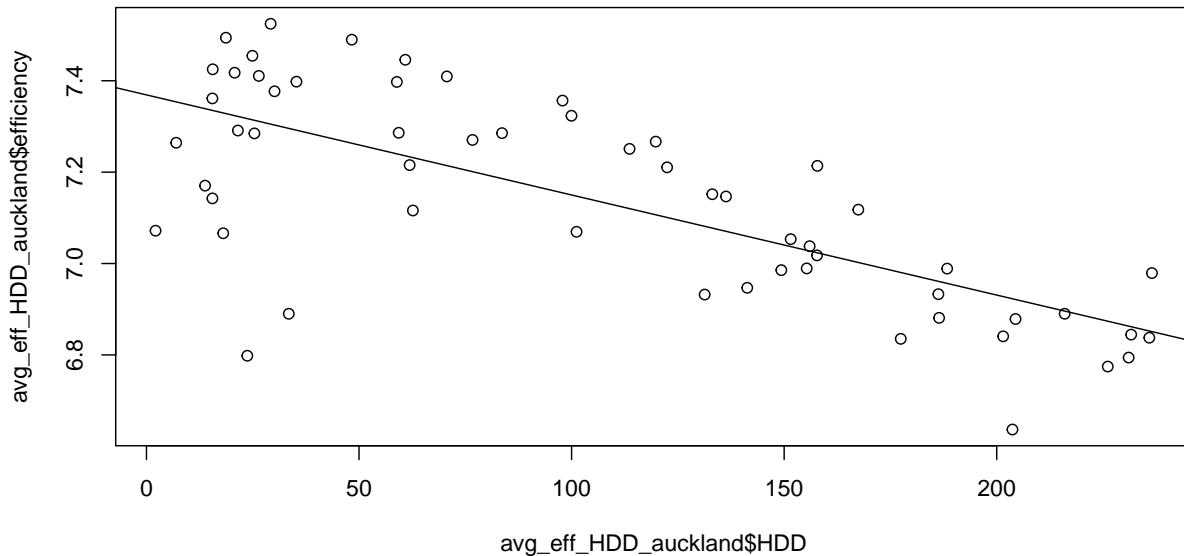
## lm(formula = efficiency ~ HDD, data = avg_eff_HDD_auckland)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.51890 -0.07959  0.01060  0.10976  0.22660
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 7.3690366  0.0362540 203.262 < 2e-16 ***
## HDD         -0.0021903  0.0002821  -7.765 2.11e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1583 on 55 degrees of freedom
## Multiple R-squared:  0.523, Adjusted R-squared:  0.5143
## F-statistic: 60.3 on 1 and 55 DF, p-value: 2.106e-10

```

```

plot(avg_eff_HDD_auckland$HDD,avg_eff_HDD_auckland$efficiency)
abline(all_lm)

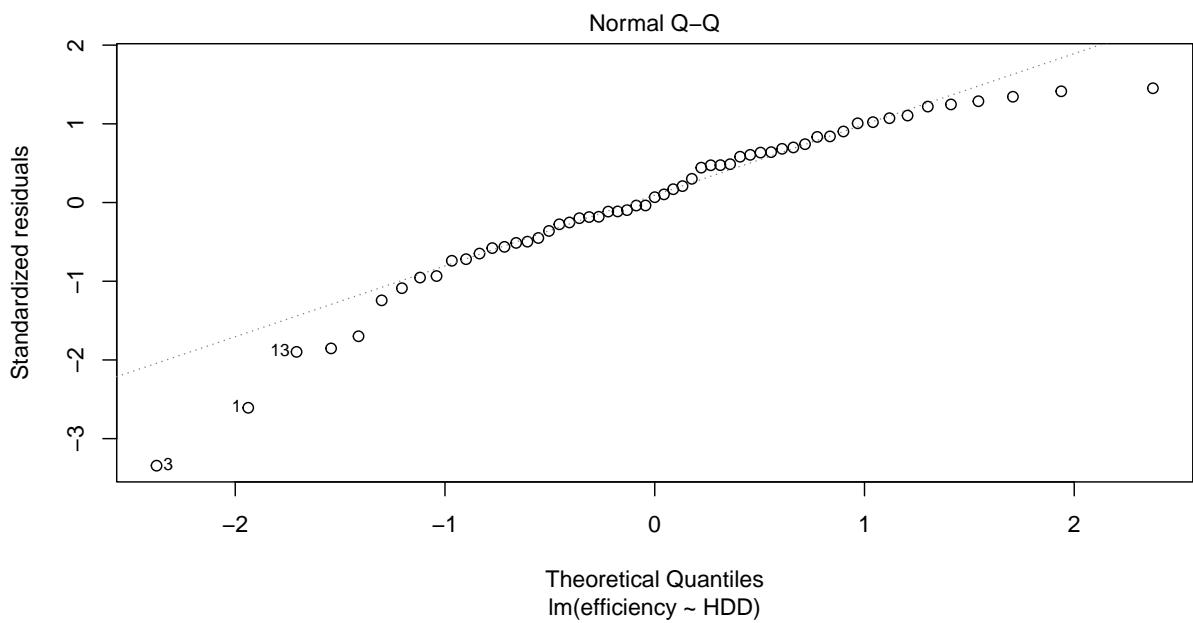
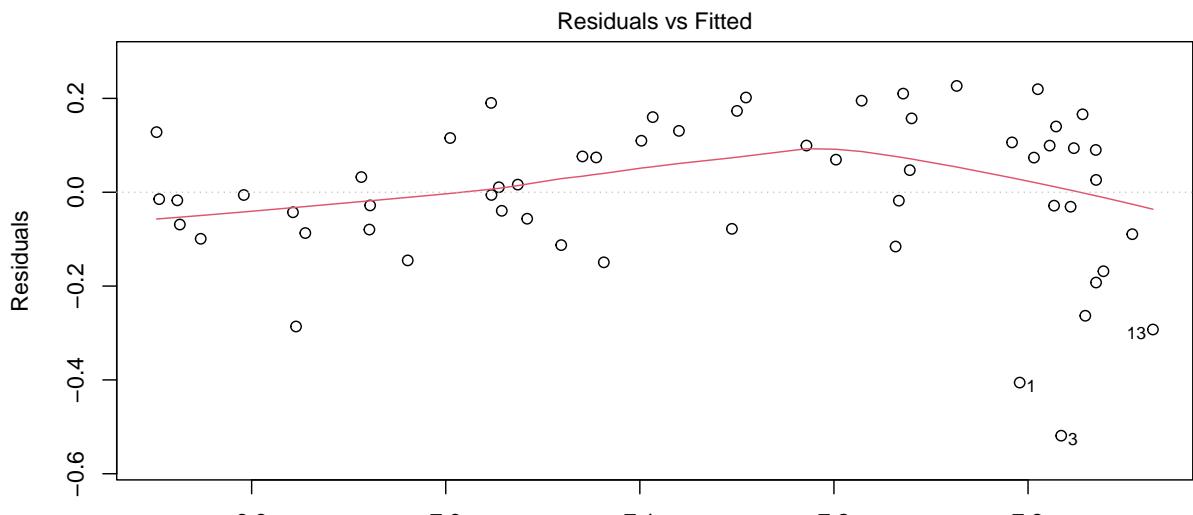
```

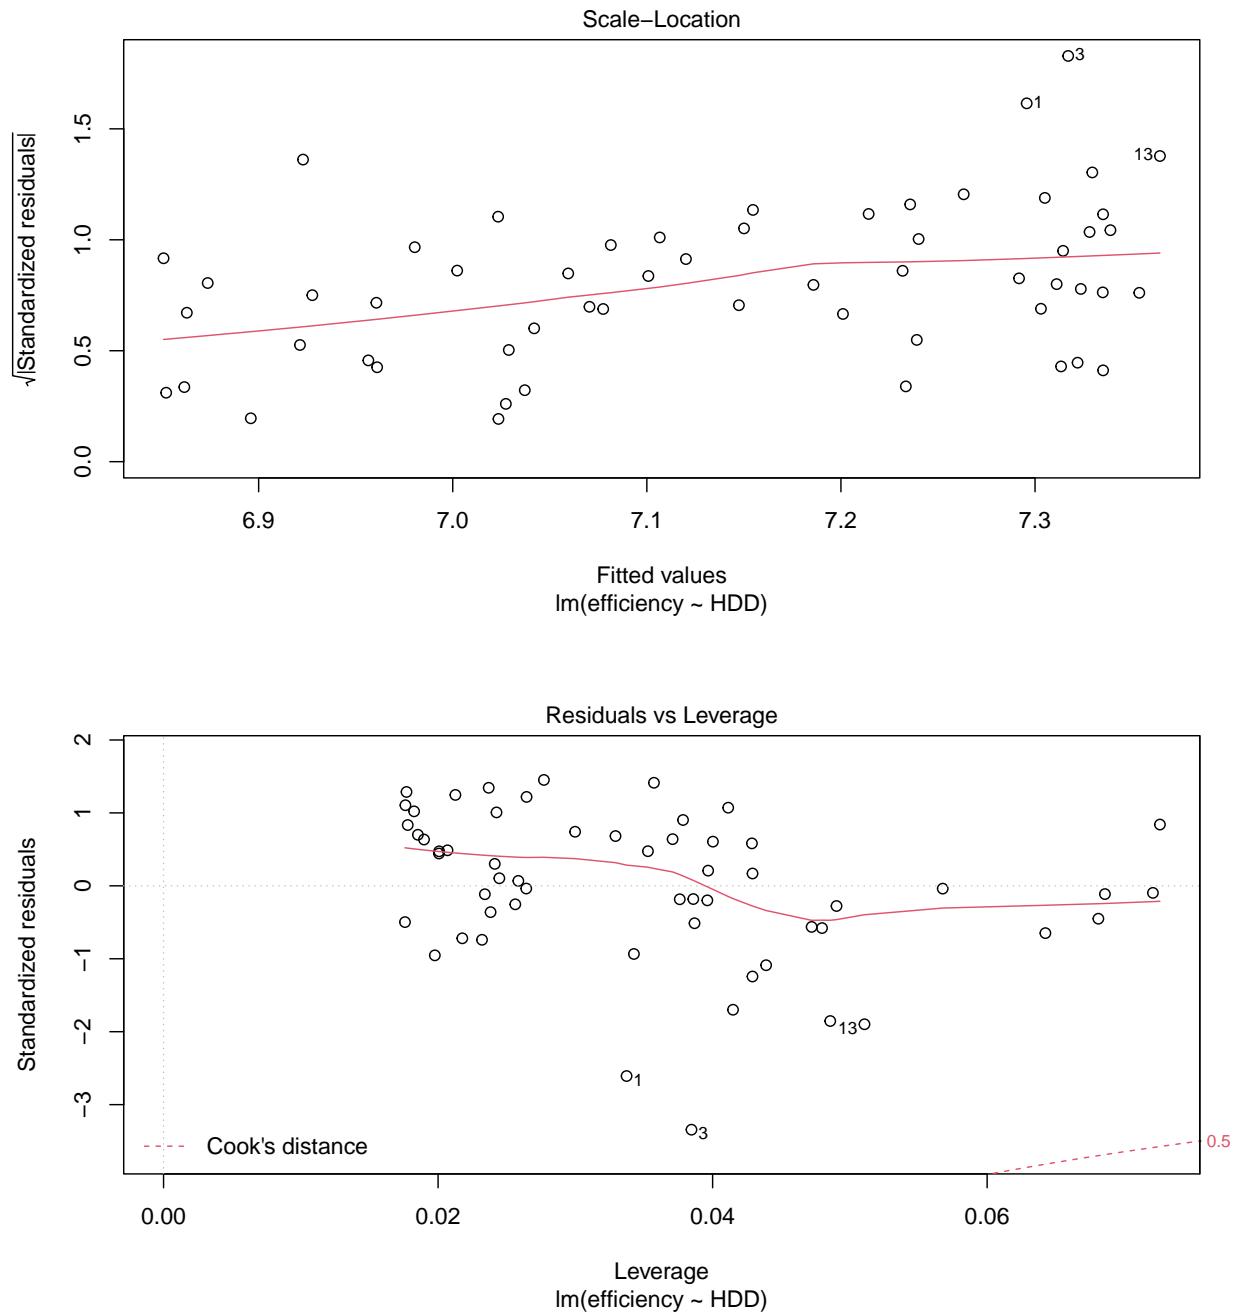


```

plot(all_lm)

```





linear model with efficiency and HDD of individual cars from 2017 to 2021 in Auckland, categorized by type of car

```
auckland_model_count = data %>%
  group_by(model) %>%
  summarise(count = n_distinct(vehicle)) %>%
  arrange(-count)
auckland_model_count
```

```
## # A tibble: 26 x 2
```

```

##      model                 count
##    <chr>                <int>
## 1 Nissan Leaf (24 kWh) 2013-2016   527
## 2 Nissan Leaf (30 kWh)           275
## 3 Nissan Leaf (24 kWh) 2011-2012   204
## 4 Nissan Leaf (40 kWh)            68
## 5 Nissan e-NV200 (24 kWh)          62
## 6 Hyundai Ioniq (EV)              26
## 7 BMW i3                         23
## 8 Hyundai Kona (EV)               14
## 9 Renault Zoe                     14
## 10 Tesla Model 3                  14
## # ... with 16 more rows

eff_HDD_auckland = data[data$year >= 2017 & region == "Auckland",]
#-2 just to check that values have in inputted correctly
eff_HDD_auckland$HDD = rep(-2, nrow(eff_HDD_auckland))

for (y in min(eff_HDD_auckland$year):max(eff_HDD_auckland$year)){
  for (m in unique(eff_HDD_auckland[eff_HDD_auckland$year == y,]$month)) {
    eff_HDD_auckland[eff_HDD_auckland$year == y & eff_HDD_auckland$month == m,]$HDD = HDD_auckland$HDD[match(y, eff_HDD_auckland$year) & match(m, eff_HDD_auckland$month)]
  }
}
eff_HDD_auckland$model = factor(eff_HDD_auckland$model, levels = auckland_model_count$model)
head(eff_HDD_auckland, 10)

##       vehicle        model region season month year distance efficiency
## 216 14840f85 Hyundai Ioniq (EV) Auckland Summer     2 2020    1371     8.5
## 217 14840f85 Hyundai Ioniq (EV) Auckland Autumn    3 2020     946     8.5
## 218 14840f85 Hyundai Ioniq (EV) Auckland Autumn    4 2020      21     8.5
## 219 14840f85 Hyundai Ioniq (EV) Auckland Autumn    5 2020     368     8.5
## 220 14840f85 Hyundai Ioniq (EV) Auckland Winter   6 2020    1109     8.3
## 221 14840f85 Hyundai Ioniq (EV) Auckland Winter   7 2020    1867     7.8
## 222 14840f85 Hyundai Ioniq (EV) Auckland Winter   8 2020     497     7.8
## 223 14840f85 Hyundai Ioniq (EV) Auckland Spring  9 2020     599     9.7
## 224 14840f85 Hyundai Ioniq (EV) Auckland Spring 10 2020    1201     9.5
## 225 14840f85 Hyundai Ioniq (EV) Auckland Spring 11 2020    1934     9.1
##      kwh      HDD
## 216 161.3 15.58750
## 217 111.3 48.30000
## 218   2.5 83.64167
## 219  43.3 136.32500
## 220 133.6 155.30833
## 221 239.4 204.42917
## 222  63.7 188.34167
## 223  61.8 157.81250
## 224 126.4  99.97917
## 225 212.5  58.90000

all_lm = lm(data = eff_HDD_auckland, efficiency~HDD)
summary(all_lm)

##

```

```

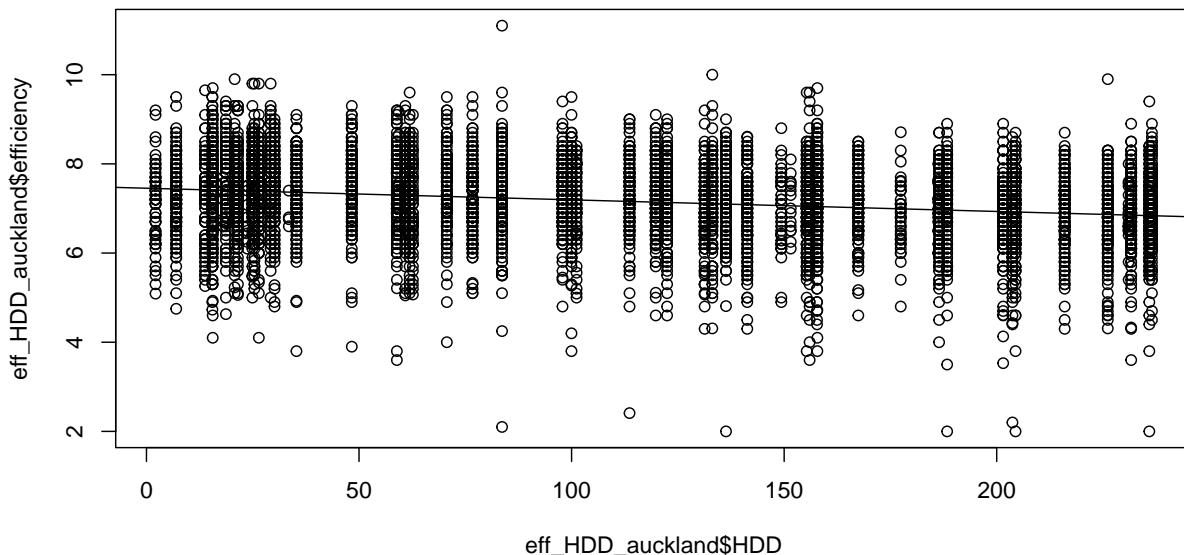
## Call:
## lm(formula = efficiency ~ HDD, data = eff_HDD_auckland)
##
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -5.1358 -0.4997  0.0424  0.5949  3.8642 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 7.4537943  0.0204472 364.54 <2e-16 ***
## HDD        -0.0026065  0.0001566 -16.65 <2e-16 ***
## ---      
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
## Residual standard error: 0.8863 on 5878 degrees of freedom
## Multiple R-squared:  0.04504,    Adjusted R-squared:  0.04487 
## F-statistic: 277.2 on 1 and 5878 DF,  p-value: < 2.2e-16

```

```

plot(eff_HDD_auckland$HDD, eff_HDD_auckland$efficiency)
abline(all_lm)

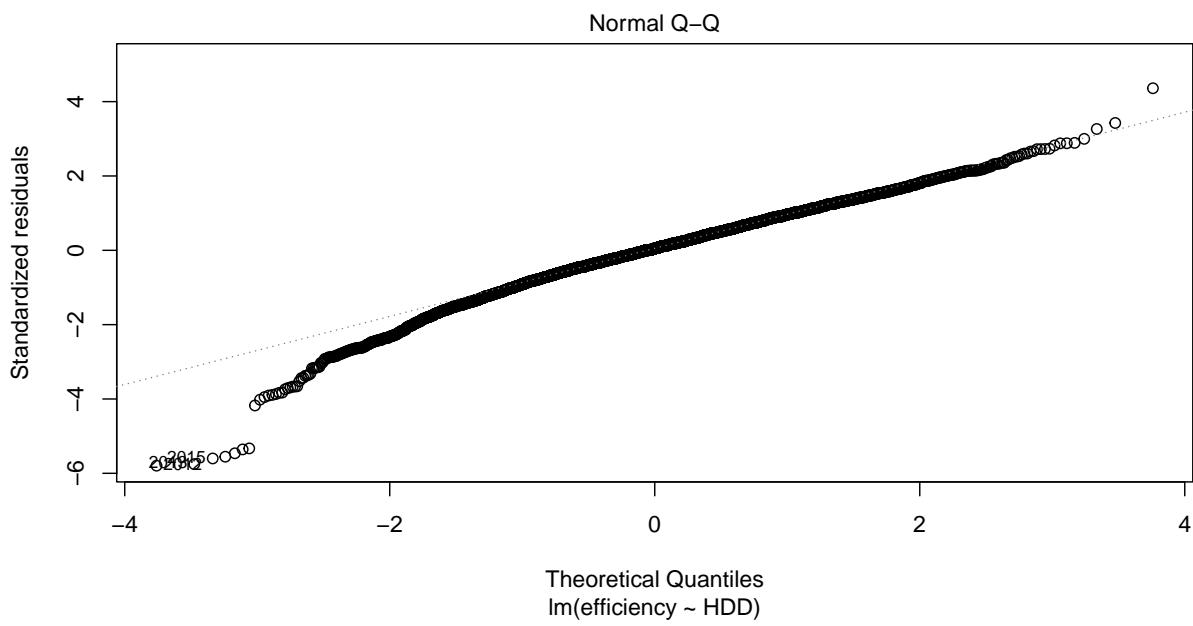
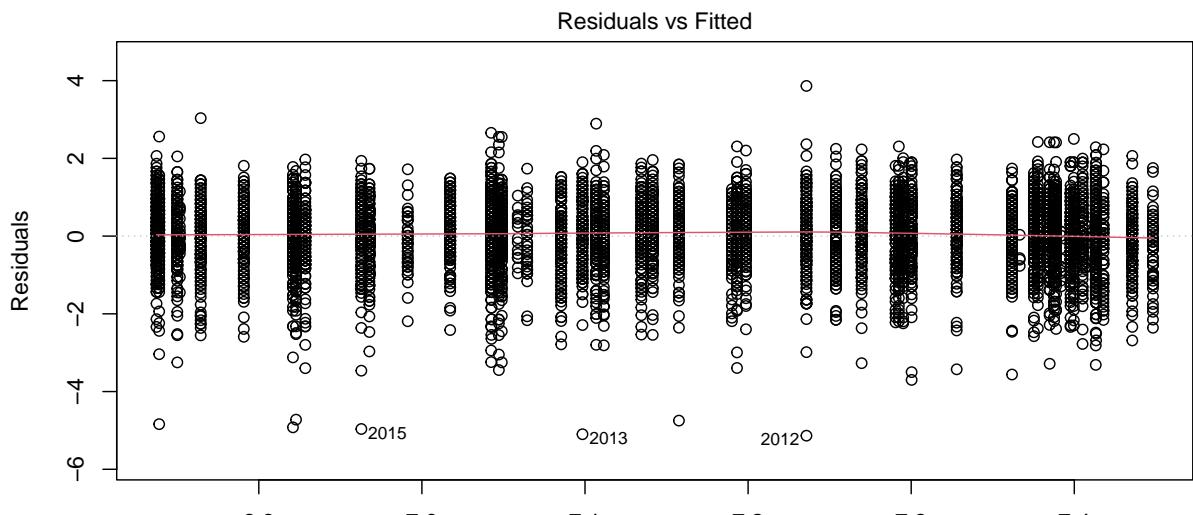
```

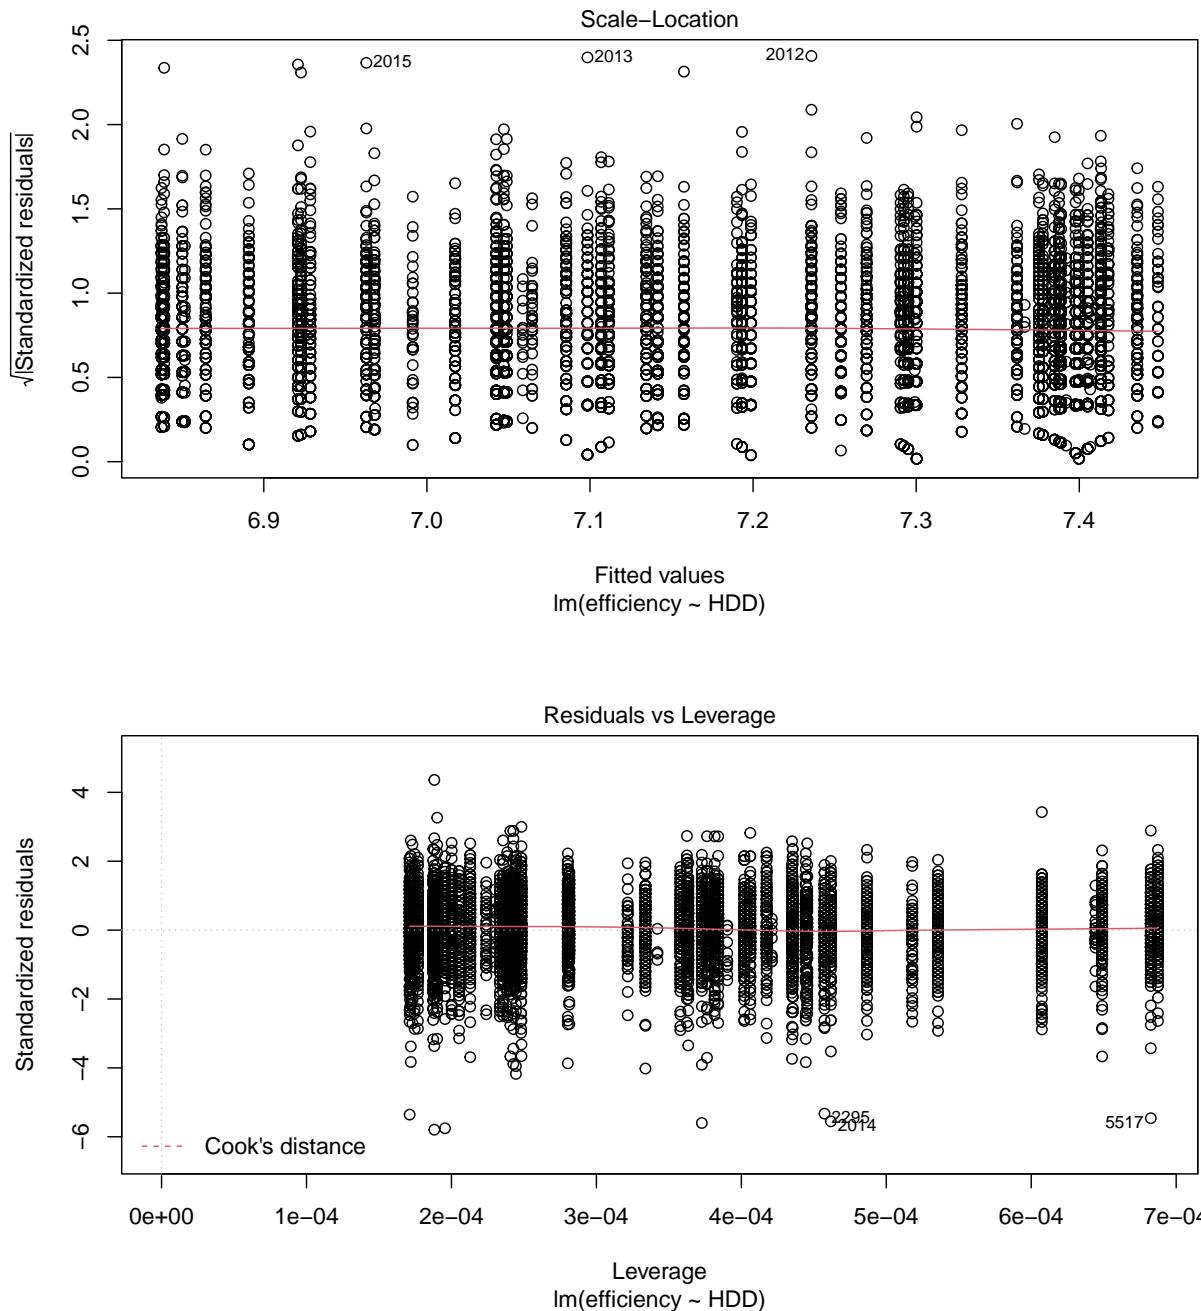


```

plot(all_lm)

```





it is using the most common EV (Nissan Leaf (24 kWh) 2013-2016) as the base line. even with many more predictors (vehicle type categories) the Std. Error of the HDD predictor decreased so t-value has actually increased

```
all_mdl_lm = lm(data = eff_HDD_auckland, efficiency~HDD+model )
summary(all_mdl_lm)
```

```
##  
## Call:
```

```

## lm(formula = efficiency ~ HDD + model, data = eff_HDD_auckland)
##
## Residuals:
##    Min     1Q Median     3Q    Max 
## -4.4941 -0.4589  0.0069  0.5004  4.9565 
##
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)    
## (Intercept)                7.784172  0.020514 379.465 < 2e-16 ***
## HDD                      -0.002549  0.000129 -19.760 < 2e-16 ***
## modelNissan Leaf (30 kWh) -0.207429  0.023246 -8.923 < 2e-16 ***
## modelNissan Leaf (24 kWh) 2011-2012 -0.976933  0.029239 -33.412 < 2e-16 ***
## modelNissan Leaf (40 kWh)   -0.680381  0.059116 -11.509 < 2e-16 ***
## modelNissan e-NV200 (24 kWh) -1.144450  0.046147 -24.800 < 2e-16 ***
## modelHyundai Ioniq (EV)      0.967094  0.116365  8.311 < 2e-16 *** 
## modelBMW i3                  -0.133892  0.084522 -1.584 0.11322  
## modelHyundai Kona (EV)       0.084274  0.096196  0.876 0.38103  
## modelRenault Zoe             -0.490255  0.111042 -4.415 1.03e-05 ***
## modelTesla Model 3           -0.685407  0.117832 -5.817 6.31e-09 *** 
## modelNissan Leaf (62 kWh)   -0.747938  0.326801 -2.289 0.02213 *  
## modelKia Niro (EV)          0.034085  0.129881  0.262 0.79300  
## modelTesla Model S           -2.401467  0.084006 -28.587 < 2e-16 *** 
## modelTesla Model-X          -3.054115  0.136349 -22.399 < 2e-16 *** 
## modelKia Soul                -0.425908  0.099548 -4.278 1.91e-05 *** 
## modelMG ZS EV                 -0.630101  0.231259 -2.725 0.00646 ** 
## modelToyota Prius            -0.417718  0.090421 -4.620 3.93e-06 *** 
## modelMitsubishi Outlander    -0.398441  0.421610 -0.945 0.34467  
## modelAudi A3 e-tron          -1.596934  0.258446 -6.179 6.89e-10 *** 
## --- 
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.7297 on 5860 degrees of freedom
## Multiple R-squared:  0.3547, Adjusted R-squared:  0.3526 
## F-statistic: 169.5 on 19 and 5860 DF,  p-value: < 2.2e-16

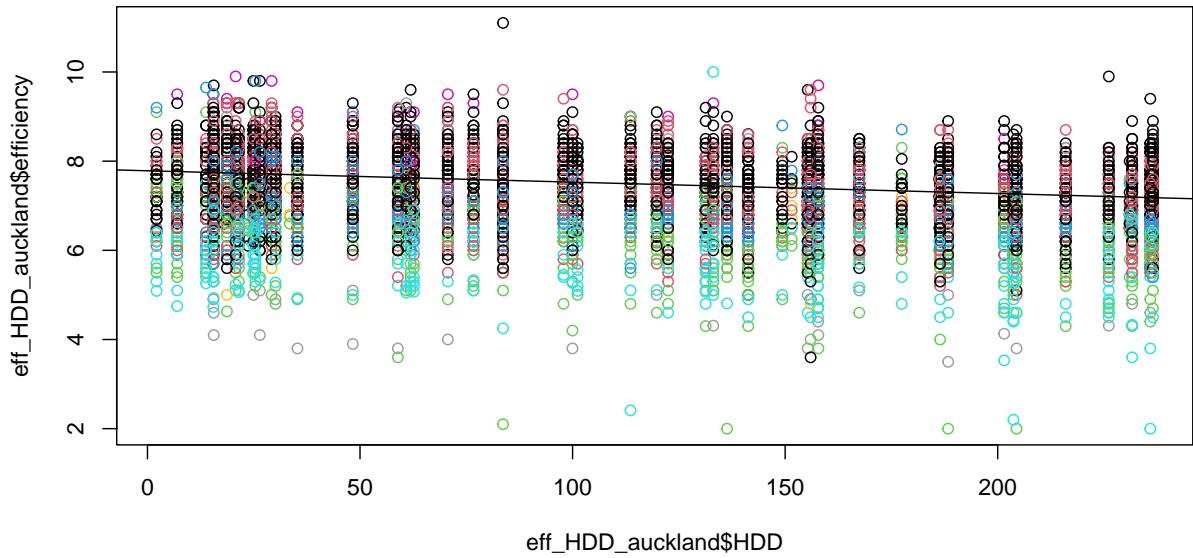
plot(eff_HDD_auckland$HDD,eff_HDD_auckland$efficiency, col = eff_HDD_auckland$model)
abline(all_mdl_lm)

```

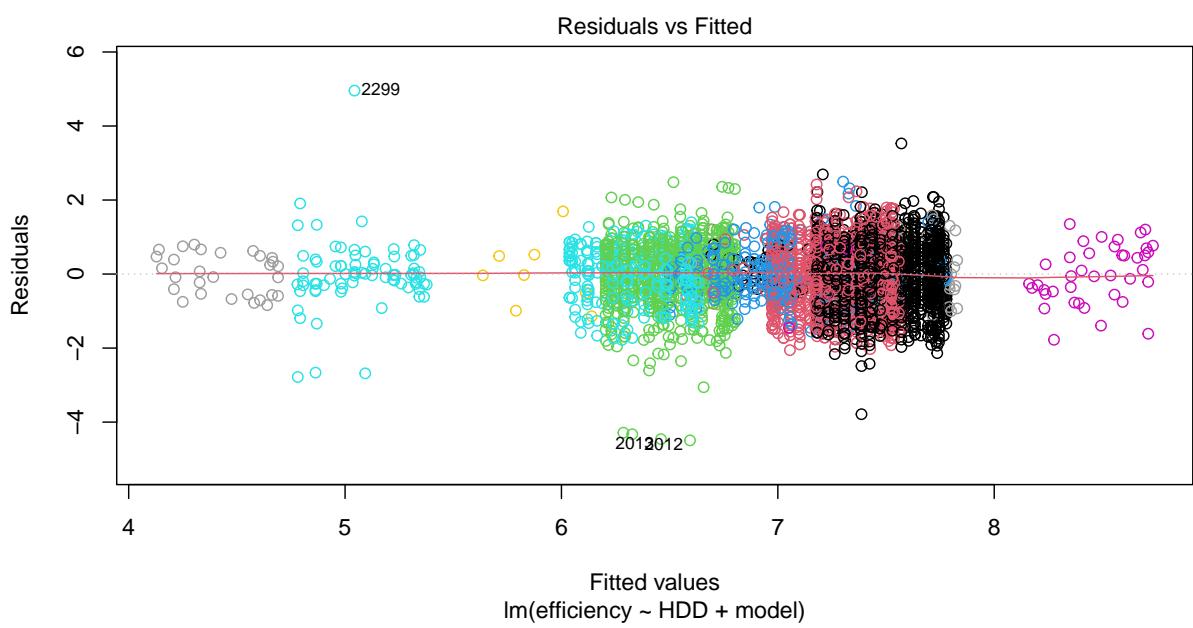
```

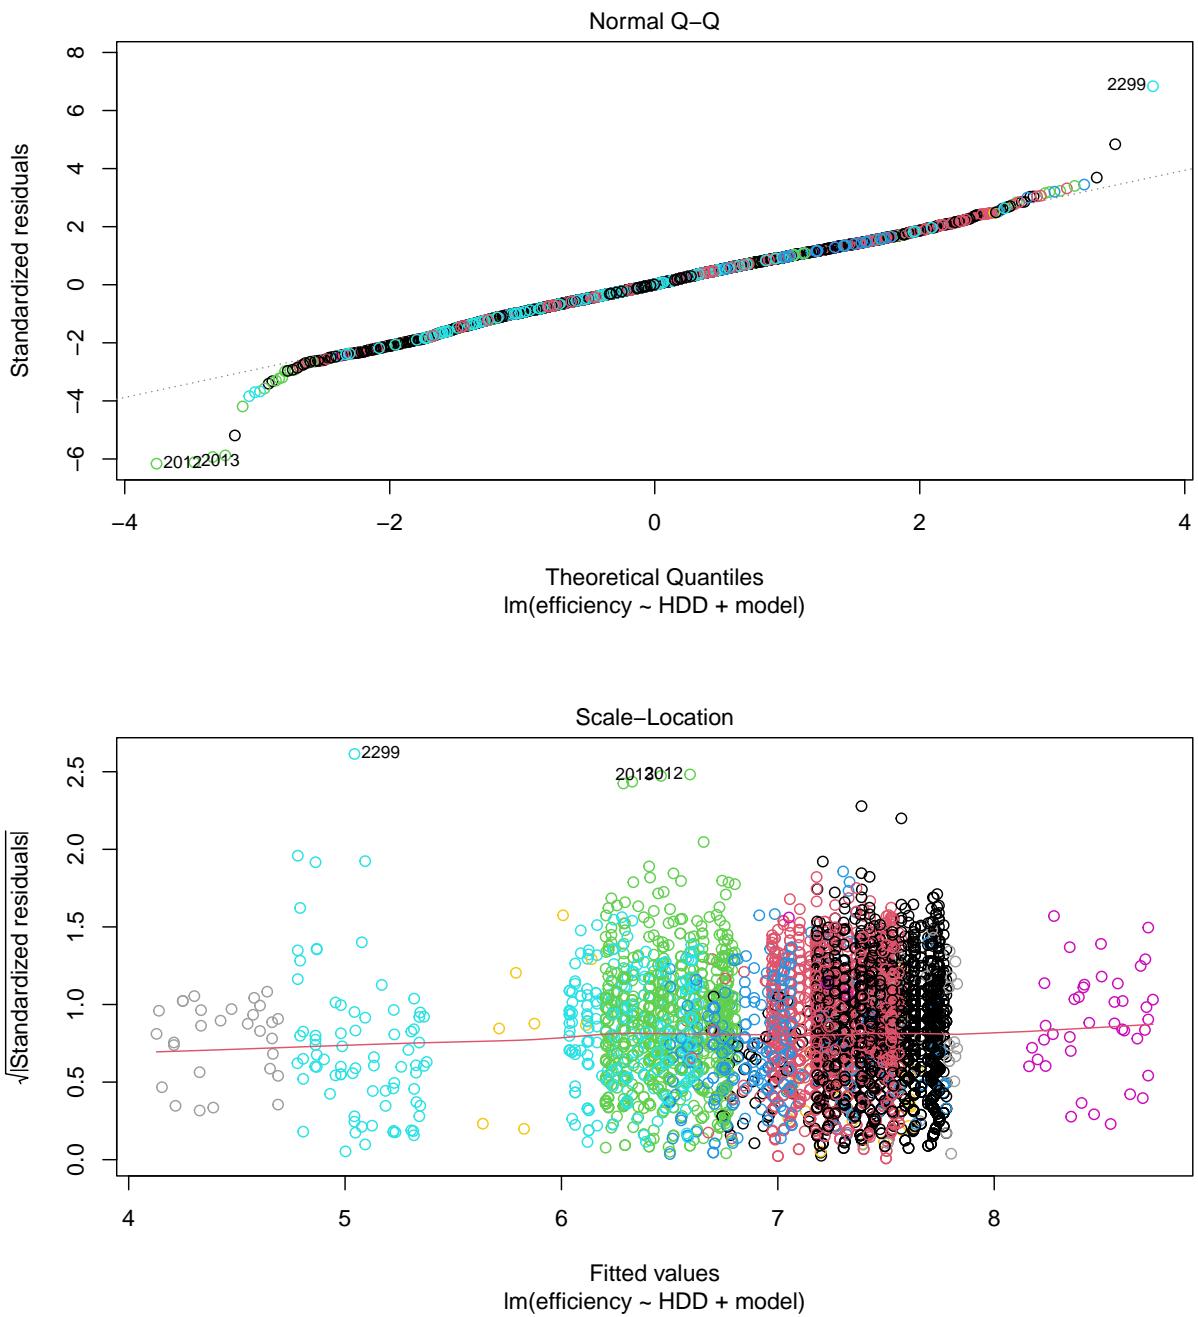
## Warning in abline(all_mdl_lm): only using the first two of 20 regression
## coefficients

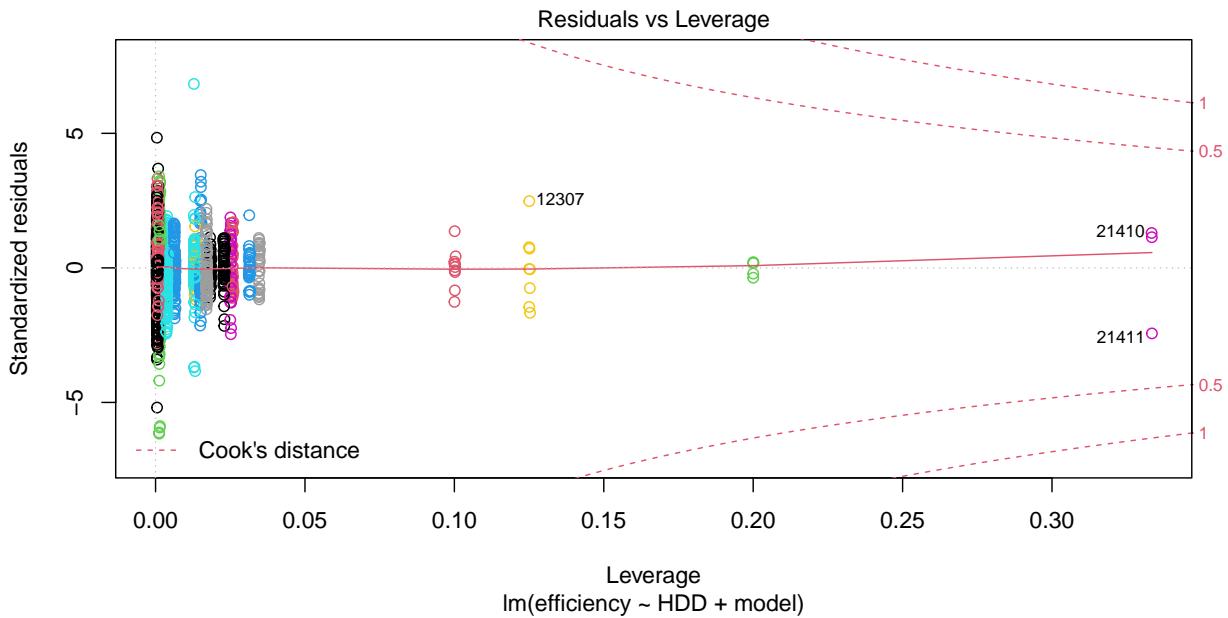
```



```
plot(all_mdl_lm, col = eff_HDD_auckland$model)
```







residual is well centered for each of the car and approx normal distribution except for a couple that have extremely low efficiency (hills, aggressive driving?)

```
inter_mdl_lm = lm(data = eff_HDD_auckland, efficiency~HDD*model )
summary(inter_mdl_lm)
```

```
##
## Call:
## lm(formula = efficiency ~ HDD * model, data = eff_HDD_auckland)
##
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -4.5202 -0.4548  0.0100  0.4969  4.9684 
##
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)    
## (Intercept)                7.7419247  0.0266540 290.460 < 2e-16  
## HDD                      -0.0021567  0.0002041 -10.567 < 2e-16  
## modelNissan Leaf (30 kWh) -0.1739057  0.0410792 -4.233 2.34e-05
## modelNissan Leaf (24 kWh)  2011-2012 -0.8133035  0.0511355 -15.905 < 2e-16 
## modelNissan Leaf (40 kWh) -0.5827516  0.1072433 -5.434 5.74e-08
## modelNissan e-NV200 (24 kWh) -1.1724003  0.0808166 -14.507 < 2e-16 
## modelHyundai Ioniq (EV)    1.4366873  0.2150252  6.681 2.58e-11
## modelBMW i3                 0.0633159  0.1486335  0.426 0.670133
## modelHyundai Kona (EV)     0.1552342  0.1802912  0.861 0.389263
## modelRenault Zoe            -0.7565049  0.1936976 -3.906 9.51e-05
## modelTesla Model 3          -0.5561063  0.2092454 -2.658 0.007890
## modelNissan Leaf (62 kWh)  -1.1472231  2.1054653 -0.545 0.585858
## modelKia Niro (EV)          0.0538757  0.2326336  0.232 0.816864
## modelTesla Model S          -2.2579304  0.1558622 -14.487 < 2e-16 
## modelTesla Model-X          -3.2253599  0.2445113 -13.191 < 2e-16 
## modelKia Soul               -0.2554847  0.1754437 -1.456 0.145385
```

```

## modelMG ZS EV           -0.9293258  0.4011377 -2.317 0.020553
## modelToyota Prius       -0.4284212  0.1533694 -2.793 0.005233
## modelMitsubishi Outlander 1.2235486  0.8151566  1.501 0.133410
## modelAudi A3 e-tron     -1.7896231  0.5189926 -3.448 0.000568
## HDD: modelNissan Leaf (30 kWh) -0.0003109  0.0003151 -0.987 0.323864
## HDD: modelNissan Leaf (24 kWh) 2011-2012 -0.0015308  0.0003927 -3.898 9.81e-05
## HDD: modelNissan Leaf (40 kWh)   -0.0009055  0.0008301 -1.091 0.275408
## HDD: modelNissan e-NV200 (24 kWh) 0.0002525  0.0006104  0.414 0.679193
## HDD: modelHyundai Ioniq (EV)    -0.0043135  0.0016606 -2.597 0.009415
## HDD: modelBMW i3            -0.0017688  0.0010886 -1.625 0.104274
## HDD: modelHyundai Kona (EV)   -0.0006516  0.0013799 -0.472 0.636780
## HDD: modelRenault Zoe        0.0024832  0.0014803  1.677 0.093501
## HDD: modelTesla Model 3      -0.0012981  0.0017980 -0.722 0.470348
## HDD: modelNissan Leaf (62 kWh) 0.0021877  0.0121543  0.180 0.857166
## HDD: modelKia Niro (EV)      -0.0001831  0.0017974 -0.102 0.918853
## HDD: modelTesla Model S      -0.0012421  0.0011060 -1.123 0.261451
## HDD: modelTesla Model-X      0.0014817  0.0017840  0.831 0.406266
## HDD: modelKia Soul          -0.0015817  0.0013419 -1.179 0.238549
## HDD: modelMG ZS EV          0.0030335  0.0032898  0.922 0.356519
## HDD: modelToyota Prius       0.0001184  0.0011941  0.099 0.921016
## HDD: modelMitsubishi Outlander -0.0220712  0.0095775 -2.304 0.021231
## HDD: modelAudi A3 e-tron     0.0016057  0.0038301  0.419 0.675071
##
## (Intercept) ***
## HDD ***
## modelNissan Leaf (30 kWh) ***
## modelNissan Leaf (24 kWh) 2011-2012 ***
## modelNissan Leaf (40 kWh) ***
## modelNissan e-NV200 (24 kWh) ***
## modelHyundai Ioniq (EV) ***
## modelBMW i3
## modelHyundai Kona (EV)
## modelRenault Zoe ***
## modelTesla Model 3 **
## modelNissan Leaf (62 kWh)
## modelKia Niro (EV)
## modelTesla Model S ***
## modelTesla Model-X ***
## modelKia Soul
## modelMG ZS EV *
## modelToyota Prius **
## modelMitsubishi Outlander
## modelAudi A3 e-tron ***
## HDD: modelNissan Leaf (30 kWh)
## HDD: modelNissan Leaf (24 kWh) 2011-2012 ***
## HDD: modelNissan Leaf (40 kWh)
## HDD: modelNissan e-NV200 (24 kWh)
## HDD: modelHyundai Ioniq (EV) **
## HDD: modelBMW i3
## HDD: modelHyundai Kona (EV)
## HDD: modelRenault Zoe .
## HDD: modelTesla Model 3
## HDD: modelNissan Leaf (62 kWh)
## HDD: modelKia Niro (EV)

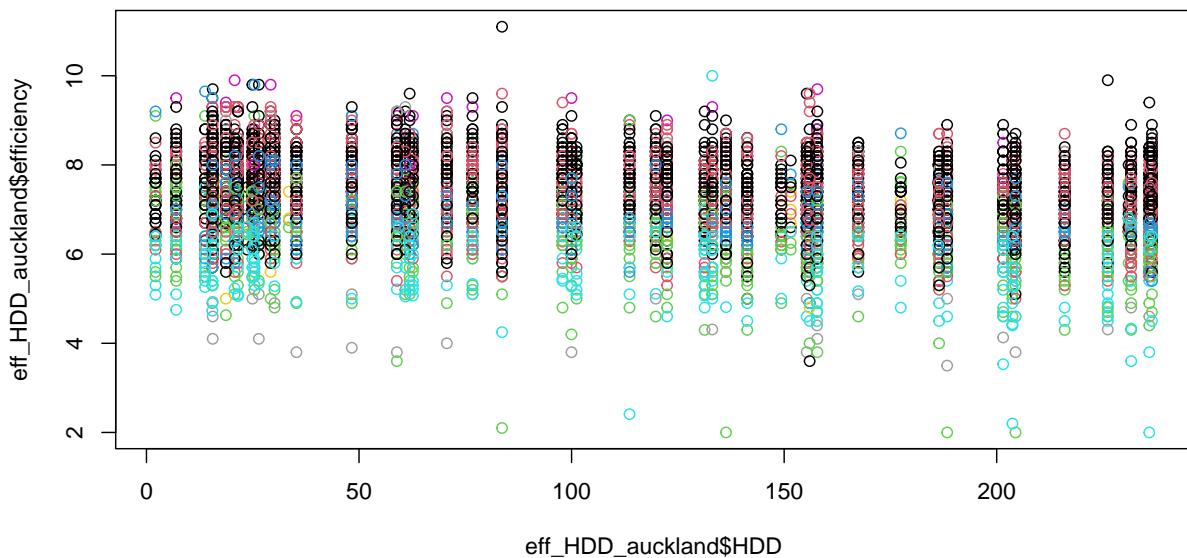
```

```

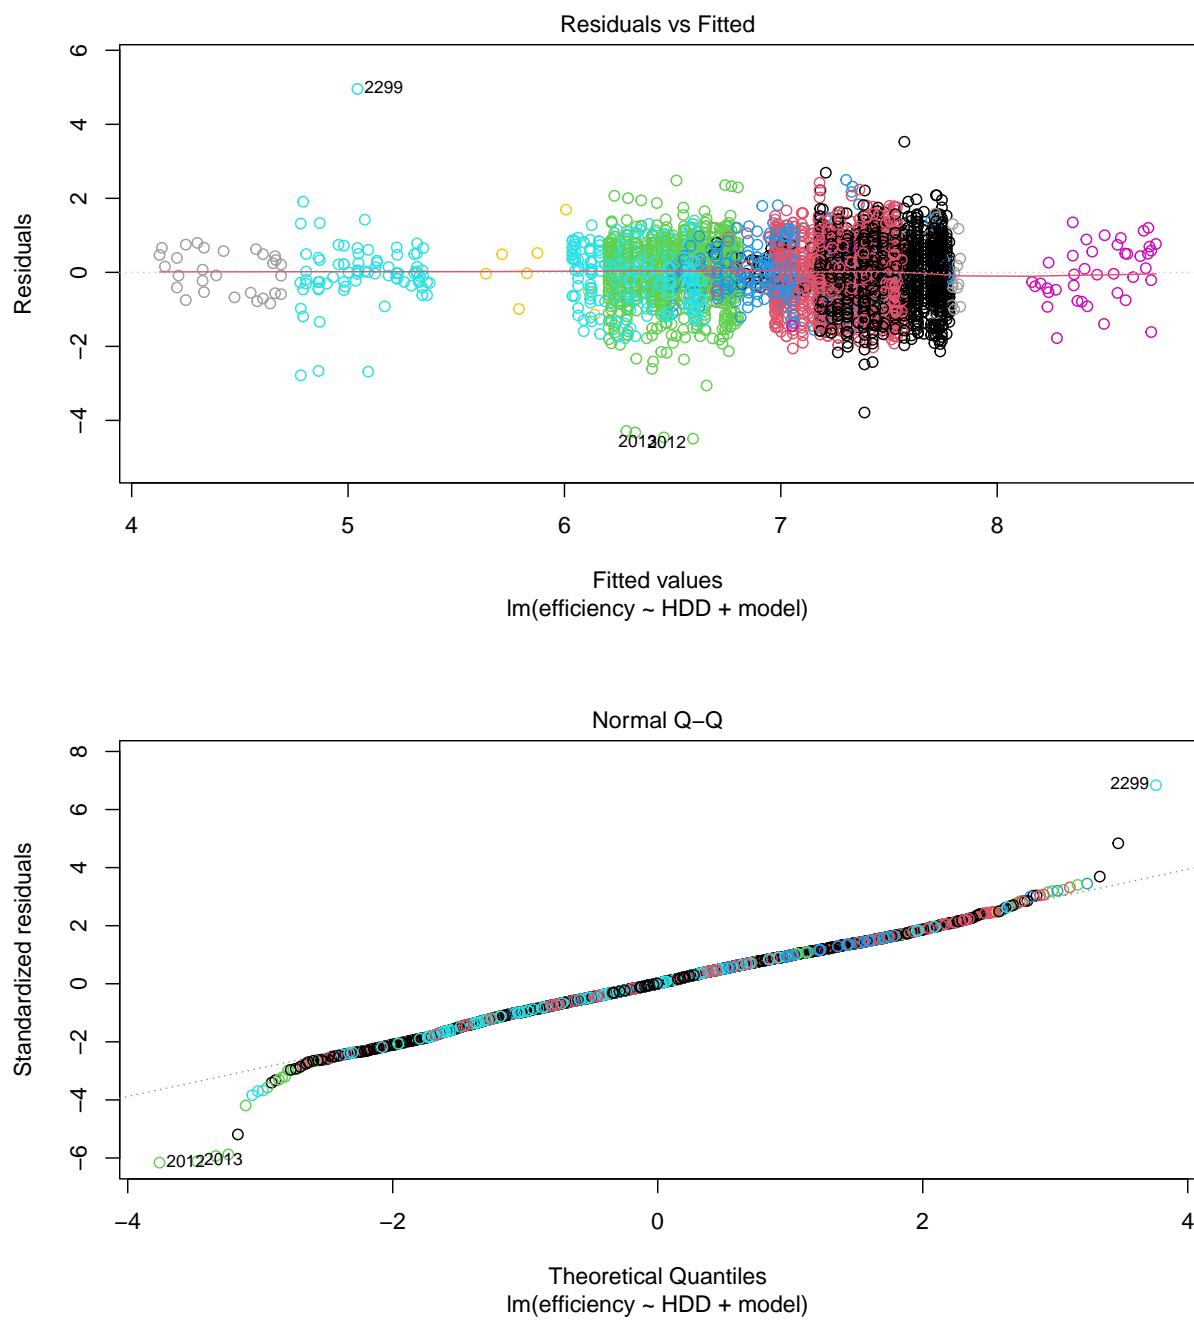
## HDD: modelTesla Model S
## HDD: modelTesla Model-X
## HDD: modelKia Soul
## HDD: modelMG ZS EV
## HDD: modelToyota Prius
## HDD: modelMitsubishi Outlander *
## HDD: modelAudi A3 e-tron
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.7285 on 5842 degrees of freedom
## Multiple R-squared: 0.3588, Adjusted R-squared: 0.3547
## F-statistic: 88.35 on 37 and 5842 DF, p-value: < 2.2e-16

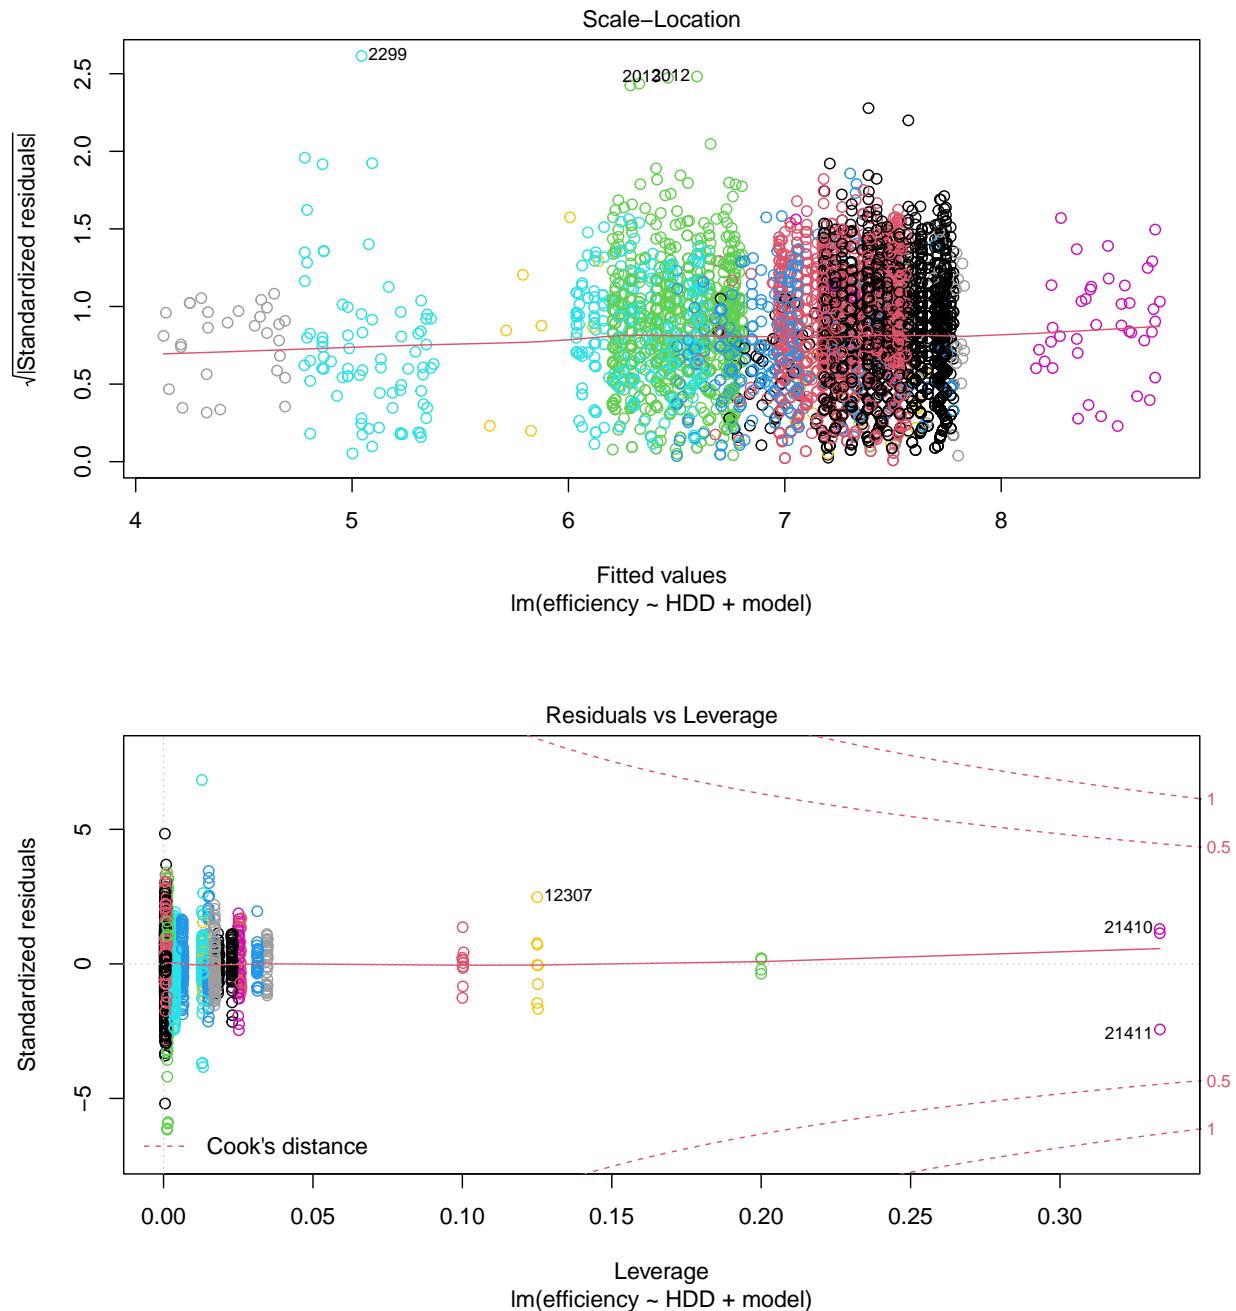
```

```
plot(eff_HDD_auckland$HDD, eff_HDD_auckland$efficiency, col = eff_HDD_auckland$model)
```



```
plot(all_mdl_lm, col = eff_HDD_auckland$model)
```





including model of car in model is worth it, interaction term anova says is significant but fit does not really improve and most are not significant. i.e not really worth being concerned about how different car respond to difference in temps

```
anova(all_lm, all_mdl_lm, inter_mdl_lm)
```

```
## Analysis of Variance Table
##
## Model 1: efficiency ~ HDD
```

```

## Model 2: efficiency ~ HDD + model
## Model 3: efficiency ~ HDD * model
##   Res.Df   RSS Df Sum of Sq      F    Pr(>F)
## 1    5878 4617.8
## 2    5860 3120.6 18   1497.24 156.7260 < 2.2e-16 ***
## 3    5842 3100.6 18     20.01  2.0951  0.004319 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```

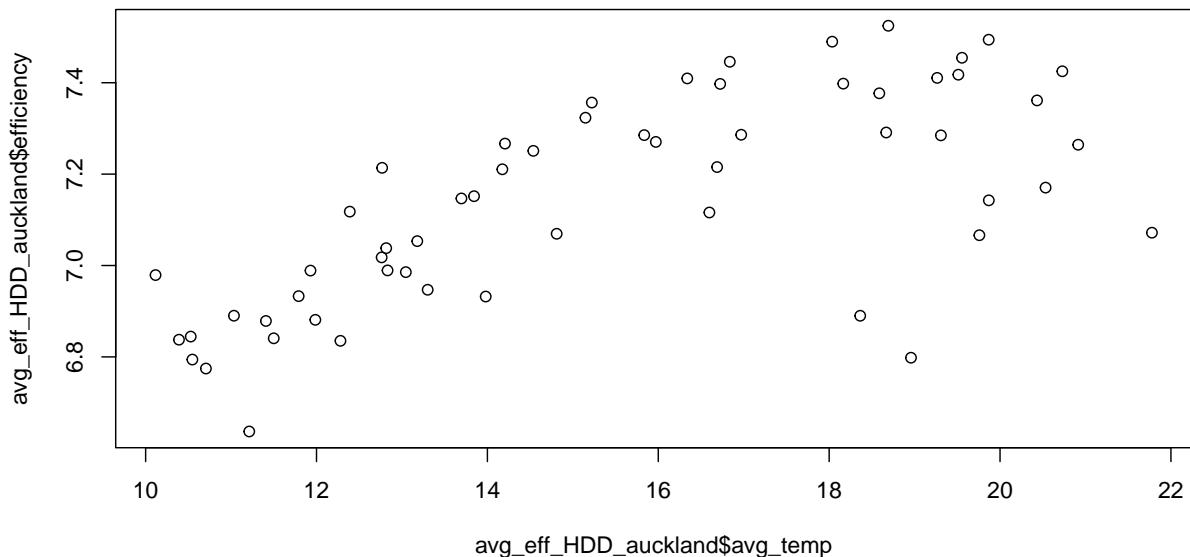
auck_avg_temp = temp_data %>%
  group_by(Year, Month) %>%
  summarise(avg_temp = mean(Temp))

```

‘summarise()’ has grouped output by ‘Year’. You can override using the ‘.groups’ argument.

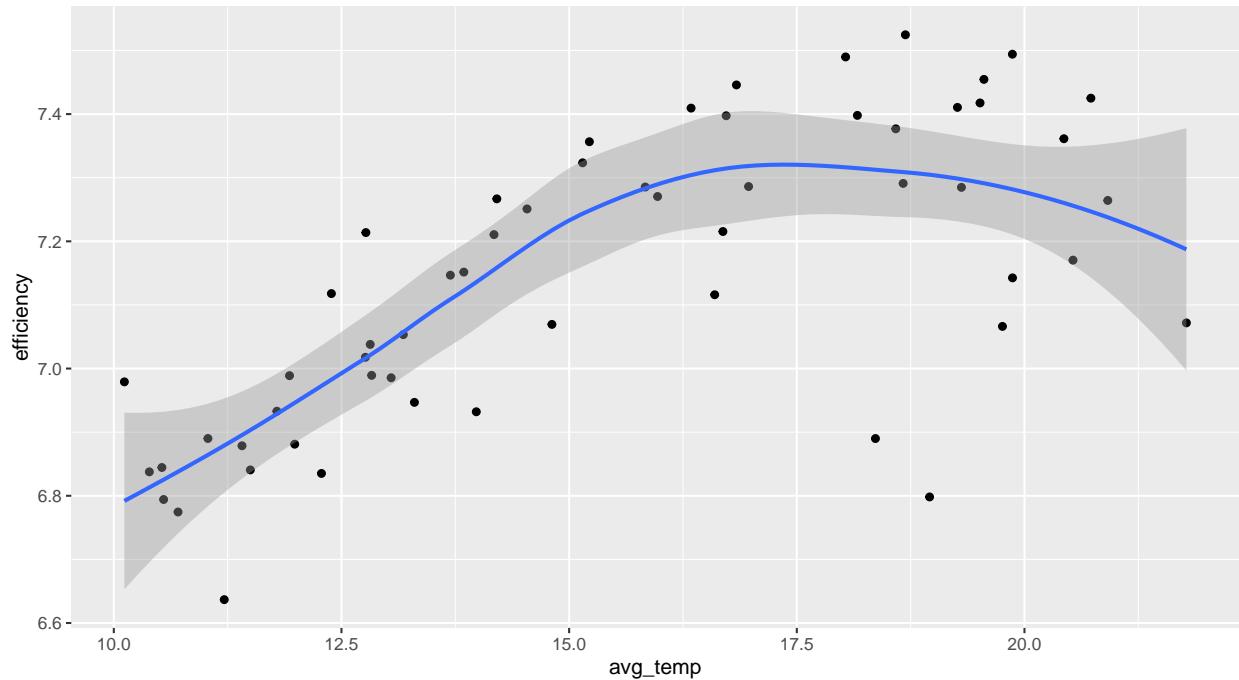
```
avg_eff_HDD_auckland = merge(avg_eff_HDD_auckland, auck_avg_temp, by.x = c('year', 'Month'), by.y = c("
```

```
plot(avg_eff_HDD_auckland$avg_temp, avg_eff_HDD_auckland$efficiency)
```



```
avg_eff_HDD_auckland %>% ggplot(aes(avg_temp, efficiency)) + geom_point() + geom_smooth(method = 'loess')
```

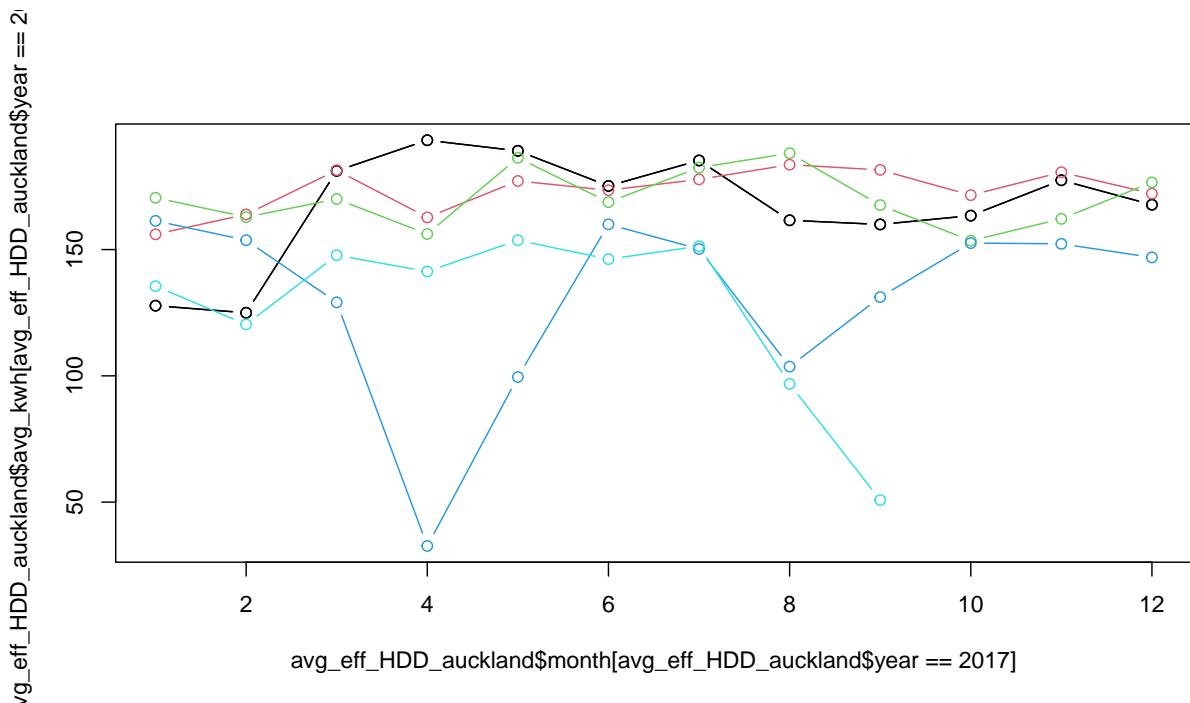
```
## ‘geom_smooth()’ using formula ‘y ~ x’
```



```

plot(avg_eff_HDD_auckland$month[avg_eff_HDD_auckland$year == 2017], avg_eff_HDD_auckland$avg_kwh[avg_eff_HDD_auckland$year == 2017], type = 'b', col = 1, ylim = c(min(avg_eff_HDD_auckland$avg_kwh),max(avg_eff_HDD_auckland$avg_kwh)))
for (i in 2017:2021) {
  points(avg_eff_HDD_auckland$month[avg_eff_HDD_auckland$year == i], avg_eff_HDD_auckland$avg_kwh[avg_eff_HDD_auckland$year == i], type = 'b', col = i-2016, xaxt = "n", ylim = c(min(avg_eff_HDD_auckland$avg_kwh),max(avg_eff_HDD_auckland$avg_kwh)))
}

```

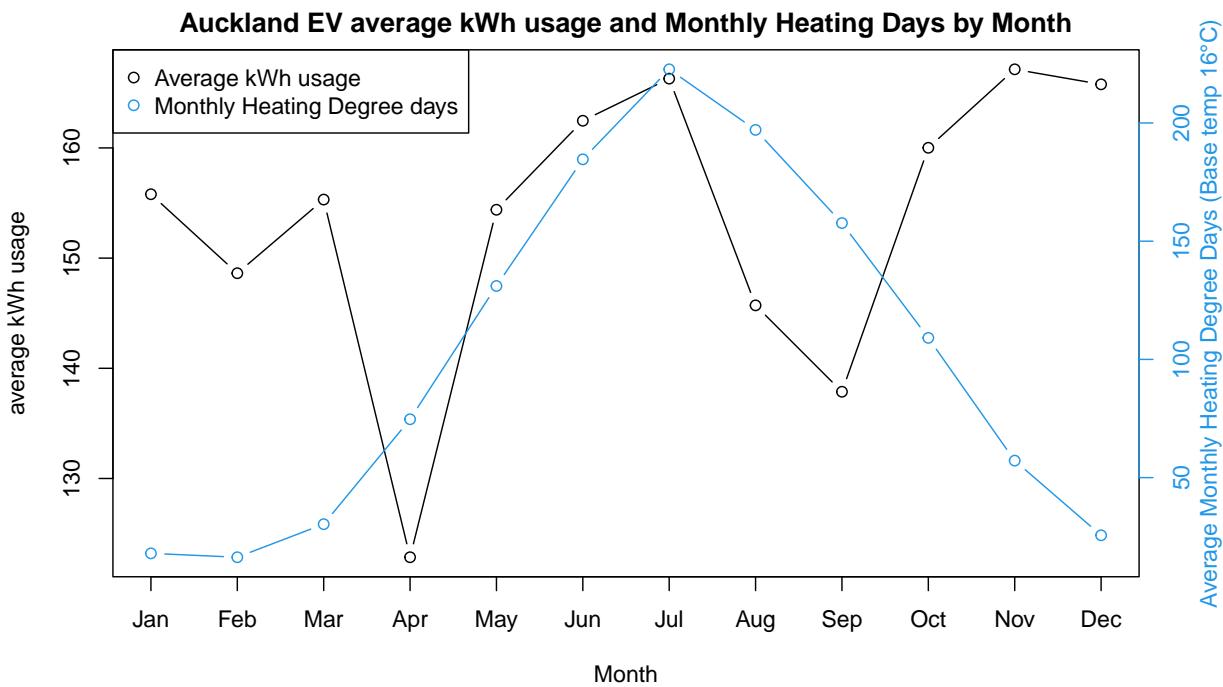


```

par(mar = c(4, 4, 2, 4))
plot(as.numeric(mean_HDD_auckland$Month), mean_HDD_auckland$kwh, type = 'b', xaxt = "n",
     ylab = "average kWh usage", xlab = "Month", main = "Auckland EV average kWh usage and Monthly Heating Degree Days by Month",
     axis(1, labels = as.character(mean_HDD_auckland$Month), at = as.numeric(mean_HDD_auckland$Month))

par(new=TRUE)
plot(as.numeric(mean_HDD_auckland$Month), mean_HDD_auckland$average_HDD, xlab="", ylab="", xlim = c(1,12),
     ylim=c(min(mean_HDD_auckland$average_HDD),max(mean_HDD_auckland$average_HDD)),
     axes=FALSE, type="b", col=4)
mtext("Average Monthly Heating Degree Days (Base temp 16°C)", side=4, col=4, line=2, cex = 1)
axis(4, ylim=c(min(mean_HDD_auckland$average_HDD),max(mean_HDD_auckland$average_HDD)),
     col=4, col.axis=4)
legend("topleft", legend = c( "Average kWh usage", "Monthly Heating Degree days"), col = c(1, 4), pch =

```



lockdowns make it too hard to see much meaningful pattern

```

mean_HDD_auckland = HDD_auckland %>%
  group_by(Month) %>%
  summarise(average_HDD = mean(HDD))

power = c()
eff = c()
kwh = c()
km = c()
for (i in 1:12) {
  eff[i] = mean(data[month==i & region == "Auckland" & year >= 2017 & year <= 2019]$efficiency, na.rm = T)
  kwh[i] = mean(data[month==i & region == "Auckland" & year >= 2017 & year <= 2019]$kwh, na.rm = T)
  km[i] = mean(data[month==i & region == "Auckland" & year >= 2017 & year <= 2019]$distance, na.rm = T)
}

mean_HDD_auckland$power_usage = power

```

```

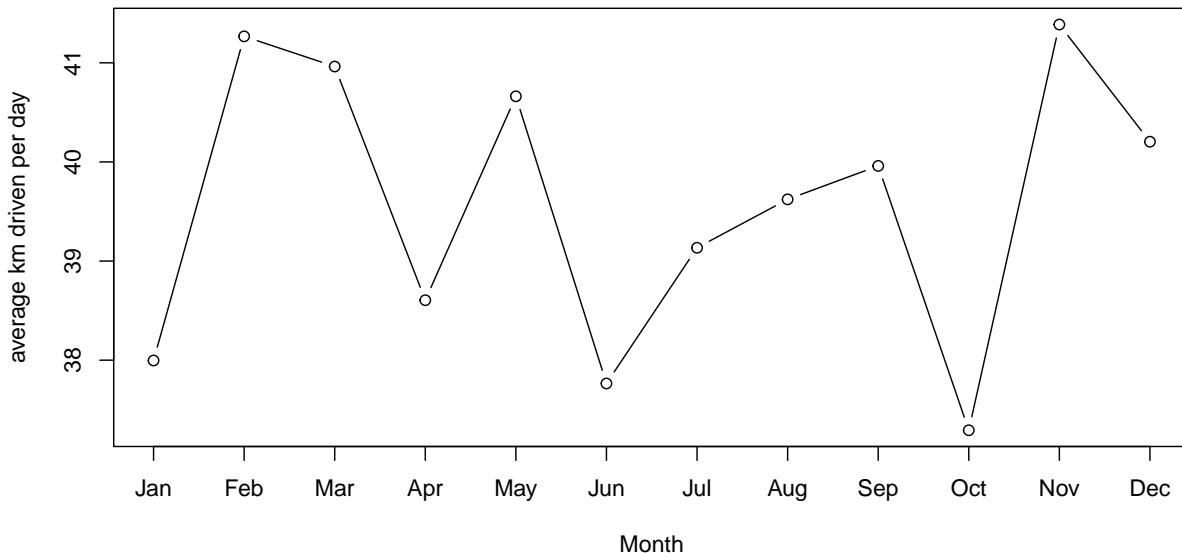
mean_HDD_auckland$efficiency = eff
mean_HDD_auckland$kwh = kwh
mean_HDD_auckland$km = km
mean_HDD_auckland

## # A tibble: 12 x 5
##   Month average_HDD efficiency     kwh     km
##   <fct>     <dbl>      <dbl> <dbl> <dbl>
## 1 Jan        18.0       7.20  165. 1178.
## 2 Feb        16.3       7.29  162. 1155.
## 3 Mar        30.3       7.36  175. 1270.
## 4 Apr         74.7       7.29  161. 1158.
## 5 May        131.       7.09  182. 1261.
## 6 Jun        185.       6.83  171. 1133.
## 7 Jul         223.       6.81  181. 1213.
## 8 Aug         197.       6.84  183. 1228.
## 9 Sep         158.       7.06  173. 1199.
## 10 Oct        109.       7.18  163. 1156.
## 11 Nov         57.2       7.31  173. 1242.
## 12 Dec        25.6       7.30  173. 1246.

plot(as.numeric(mean_HDD_auckland$Month), mean_HDD_auckland$km/month_length, type = 'b', xaxt = "n",
     ylab = "average km driven per day", xlab = "Month", main = "Auckland EV km driven 2017-2019")
axis(1, labels = as.character(mean_HDD_auckland$Month), at = as.numeric(mean_HDD_auckland$Month))

```

Auckland EV km driven 2017–2019



```

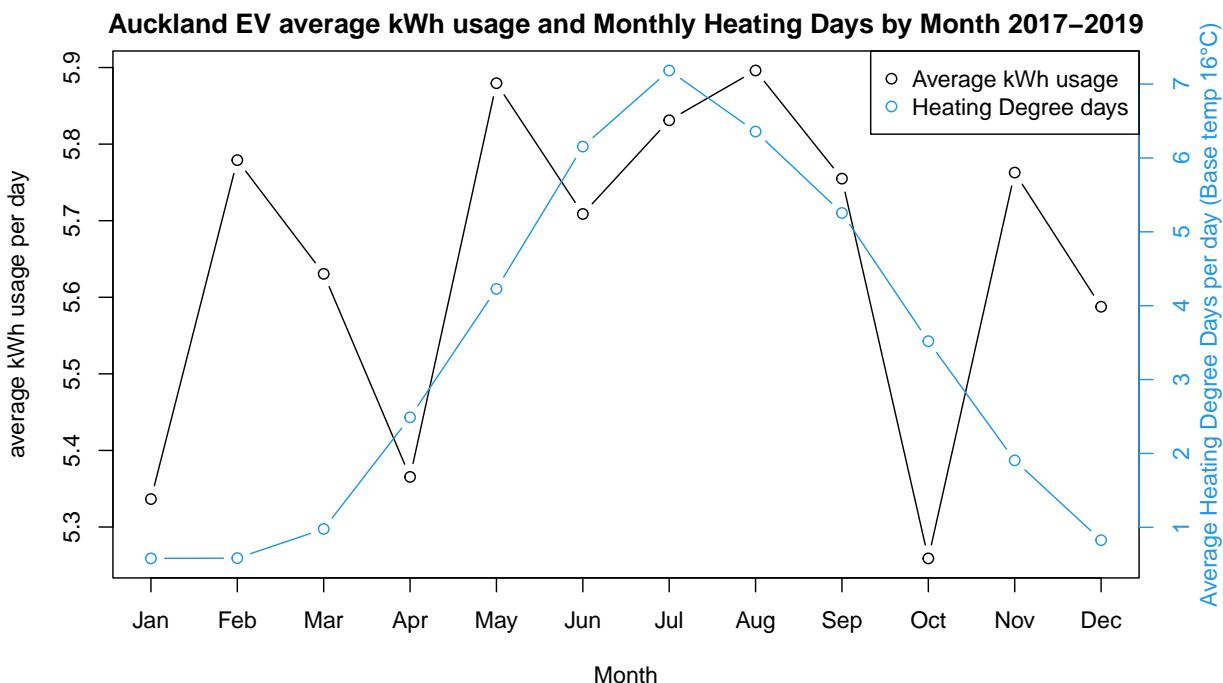
par(mar = c(4, 4, 2, 4))
plot(as.numeric(mean_HDD_auckland$Month), mean_HDD_auckland$kwh/month_length, type = 'b', xaxt = "n",
     ylab = "average kWh usage per day", xlab = "Month", main = "Auckland EV average kWh usage and Month")
axis(1, labels = as.character(mean_HDD_auckland$Month), at = as.numeric(mean_HDD_auckland$Month))

```

```

par(new=TRUE)
plot(as.numeric(mean_HDD_auckland$Month) , mean_HDD_auckland$average_HDD/month_length, xlab="", ylab="",
     ylim=c(min(mean_HDD_auckland$average_HDD/month_length),max(mean_HDD_auckland$average_HDD/month_length)),
     axes=FALSE, type="b", col=4)
mtext("Average Heating Degree Days per day (Base temp 16°C)", side=4, col=4, line=2, cex = 1)
axis(4, ylim=c(min(mean_HDD_auckland$average_HDD),max(mean_HDD_auckland$average_HDD)),
     col=4, col.axis=4)
legend("topright", legend = c( "Average kWh usage", "Heating Degree days"), col = c(1, 4), pch = 1)

```



EV power used data goes to the contrary of the fuel usage data. while it is clear that efficiency decreases in the winter months going by fuel usage data total power used should be less in the winter as there is significantly more fuel usage in the summer. is this cause petrol fuel efficiency is lower in summer (possibly due to AC), current EV drivers are different (early adopters or EV commute with 2nd petrol road trip car), or a regional thing?

```

mean_HDD_not_auckland = HDD_auckland %>%
  group_by(Month) %>%
  summarise(average_HDD = mean(HDD))

power = c()
eff = c()
kwh = c()
for (i in 1:12) {
  eff[i] = mean(data[month==i & region != "Auckland" & year >= 2017 & year <= 2019]$efficiency, na.rm = T)
  kwh[i] = mean(data[month==i & region != "Auckland" & year >= 2017 & year <= 2019]$kwh, na.rm = T)
}

mean_HDD_not_auckland$power_usage = power
mean_HDD_not_auckland$efficiency = eff
mean_HDD_not_auckland$kwh = kwh

```

```

mean_HDD_not_auckland

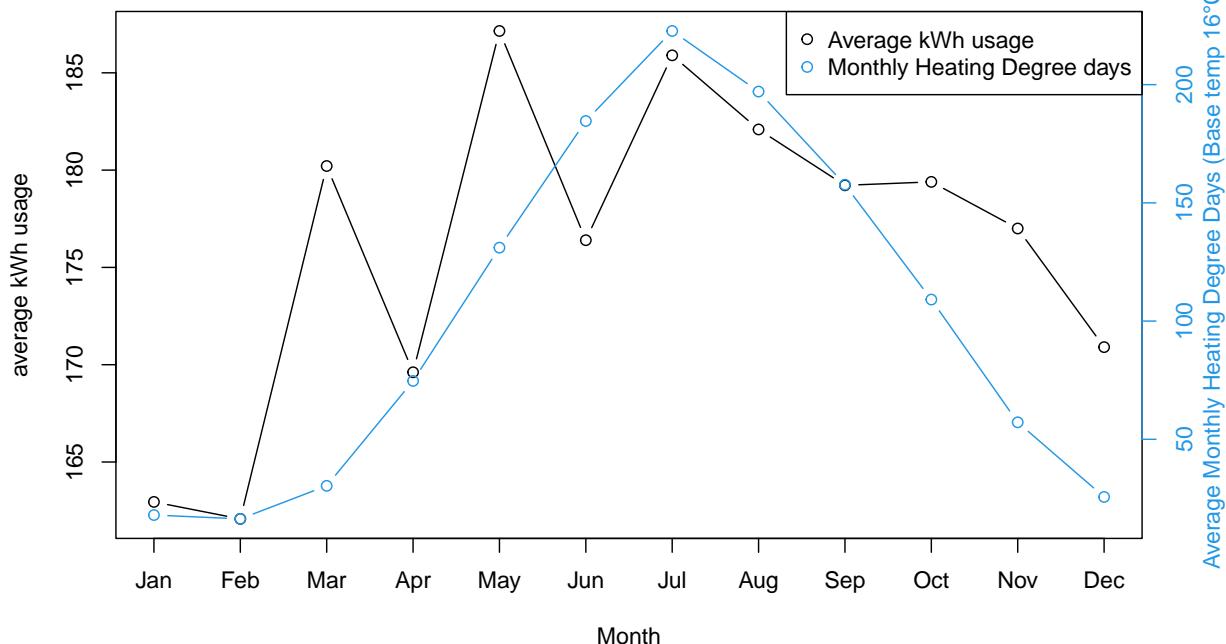
## # A tibble: 12 x 4
##   Month average_HDD efficiency     kwh
##   <fct>      <dbl>      <dbl> <dbl>
## 1 Jan        18.0       7.05  163.
## 2 Feb        16.3       7.06  162.
## 3 Mar        30.3       7.05  180.
## 4 Apr         74.7       6.77  170.
## 5 May        131.       6.64  187.
## 6 Jun        185.       6.36  176.
## 7 Jul        223.       6.32  186.
## 8 Aug        197.       6.48  182.
## 9 Sep        158.       6.69  179.
## 10 Oct       109.       6.83  179.
## 11 Nov        57.2       6.99  177
## 12 Dec       25.6       7.02  171.

#not a fair plot as is using Auckland HDD against not auckland kwh
#just used to get an idea for power usage outside auckland
par(mar = c(4, 4, 2, 4))
plot(as.numeric(mean_HDD_not_auckland$Month), mean_HDD_not_auckland$kwh, type = 'b', xaxt = "n",
     ylab = "average kWh usage", xlab = "Month", main = "not Auckland EV average kWh usage and Monthly HDD")
axis(1, labels = as.character(mean_HDD_not_auckland$Month), at = as.numeric(mean_HDD_not_auckland$Month))

par(new=TRUE)
plot(as.numeric(mean_HDD_not_auckland$Month), mean_HDD_not_auckland$average_HDD, xlab="", ylab="", xlim =
     ylim=c(min(mean_HDD_not_auckland$average_HDD),max(mean_HDD_not_auckland$average_HDD)),
     axes=FALSE, type="b", col=4)
mtext("Average Monthly Heating Degree Days (Base temp 16°C)", side=4, col=4, line=2, cex = 1)
axis(4, ylim=c(min(mean_HDD_not_auckland$average_HDD),max(mean_HDD_not_auckland$average_HDD)),
     col=4, col.axis=4)
legend("topright", legend = c( "Average kWh usage", "Monthly Heating Degree days"), col = c(1, 4), pch =

```

not Auckland EV average kWh usage and Monthly Heating Days by Month 2017–2019



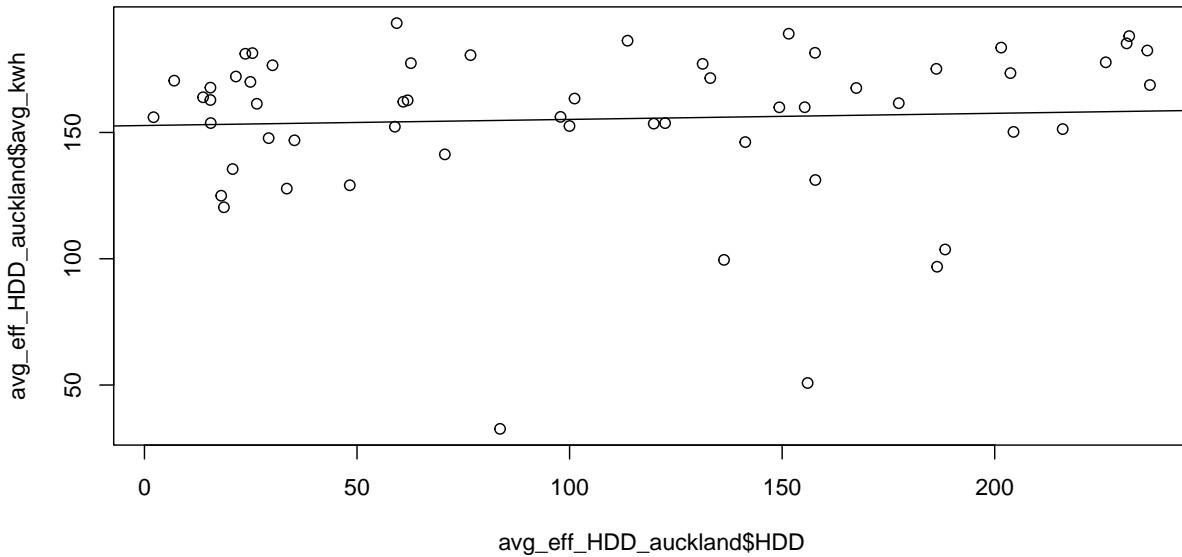
excluding Auckland from the data it appears that it is not a regional thing therefore could be petrol fuel efficiency is lower in summer (possibly) or current EV drivers are different (early adopters or EV commute with 2nd petrol road trip car)? <https://www.sciencedirect.com/science/article/pii/S1361920914000923> shows petrol fuel efficiency is lower in the winter

however <https://www.proquest.com/docview/2131026877?pq-origsite=gscholar&fromopenview=true> cites paper <https://www.sae.org/publications/technical-papers/content/2002-01-1957/> that AC accounts for 7% extra power usage

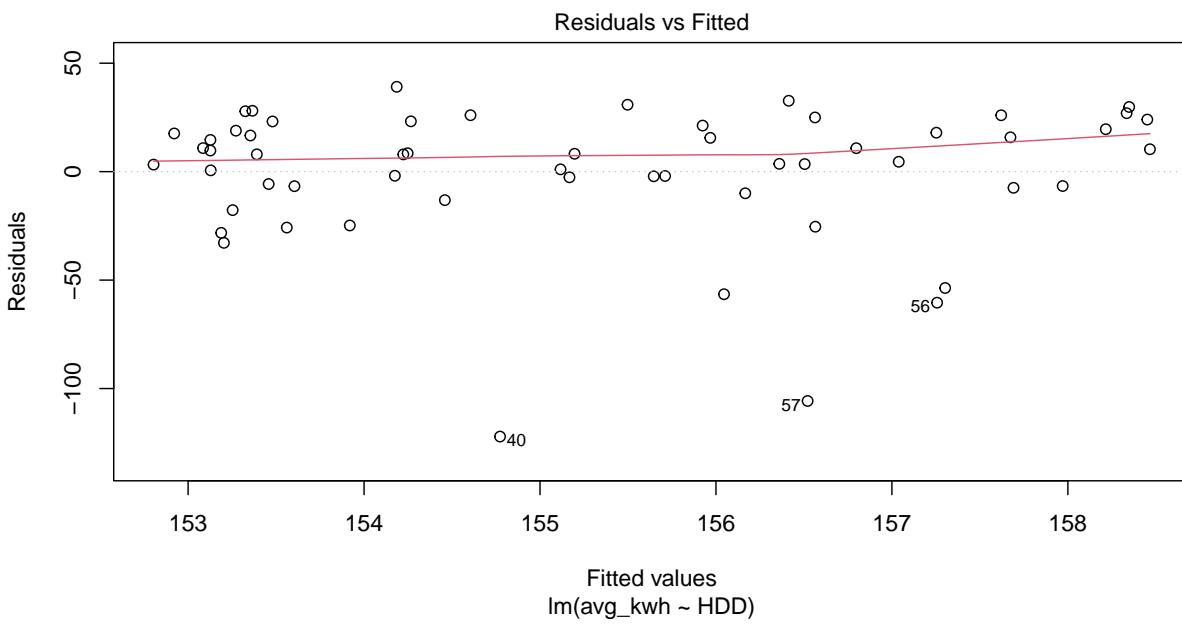
```
all_lm = lm(data = avg_eff_HDD_auckland, avg_kwh ~ HDD)
summary(all_lm)
```

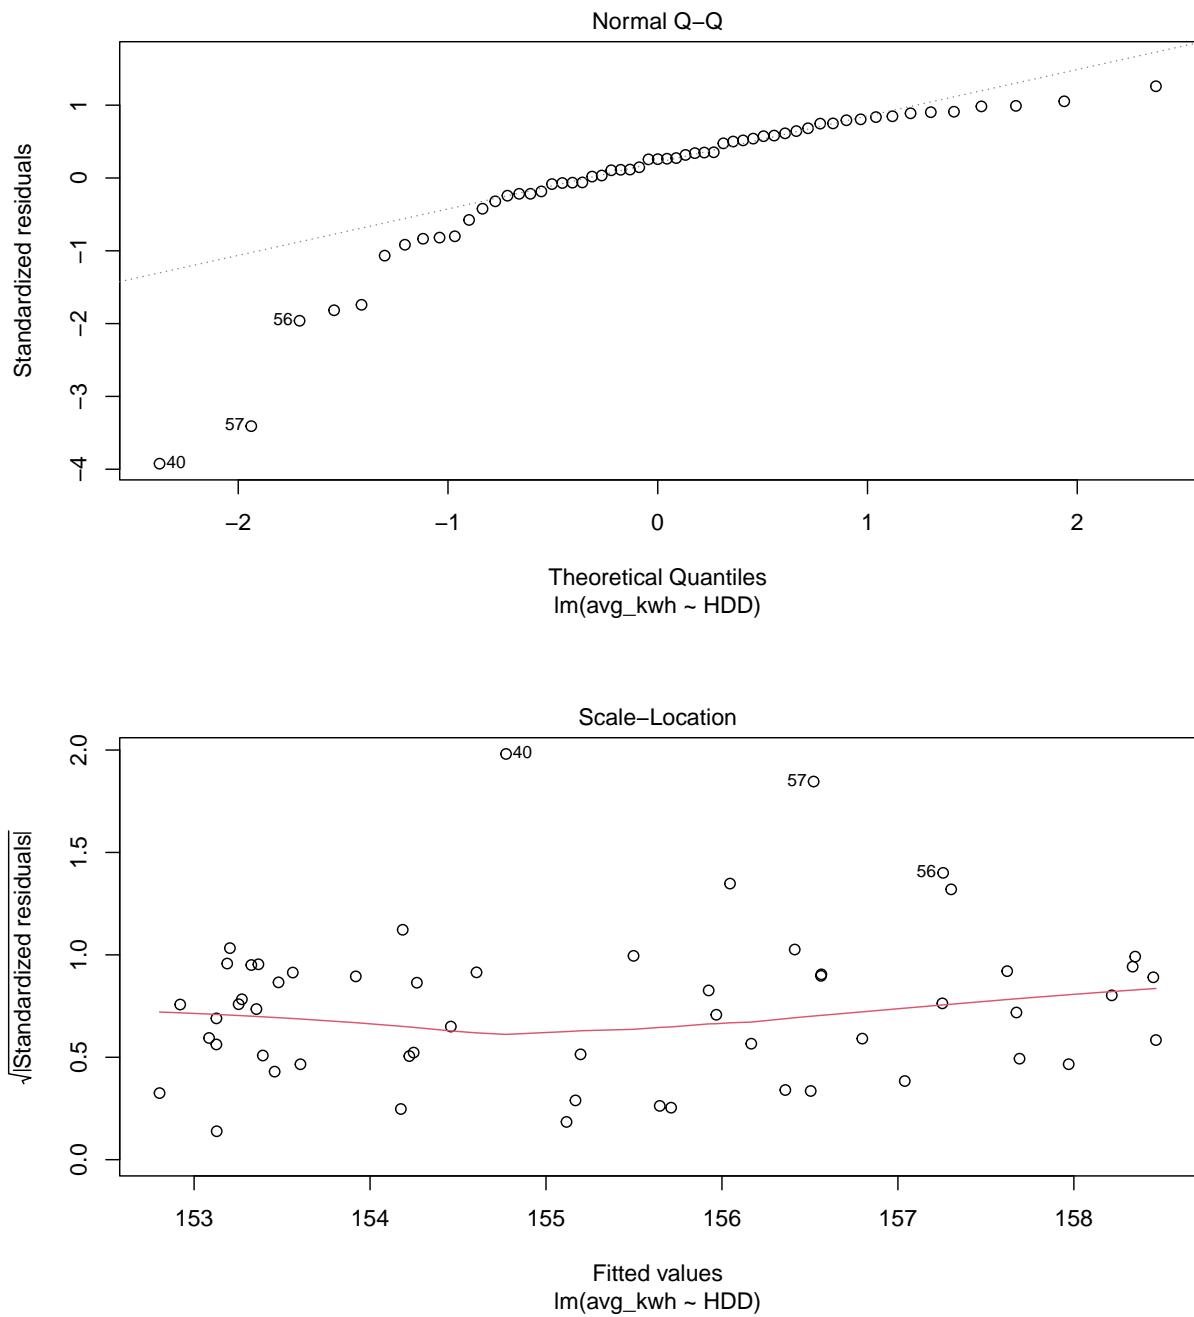
```
##
## Call:
## lm(formula = avg_kwh ~ HDD, data = avg_eff_HDD_auckland)
##
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -122.135   -6.708    7.976   19.572   39.114 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 152.75192    7.19884  21.219 <2e-16 ***
## HDD          0.02416    0.05601   0.431   0.668    
## ---        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
## Residual standard error: 31.43 on 55 degrees of freedom
## Multiple R-squared:  0.003372, Adjusted R-squared:  -0.01475 
## F-statistic: 0.1861 on 1 and 55 DF,  p-value: 0.6679
```

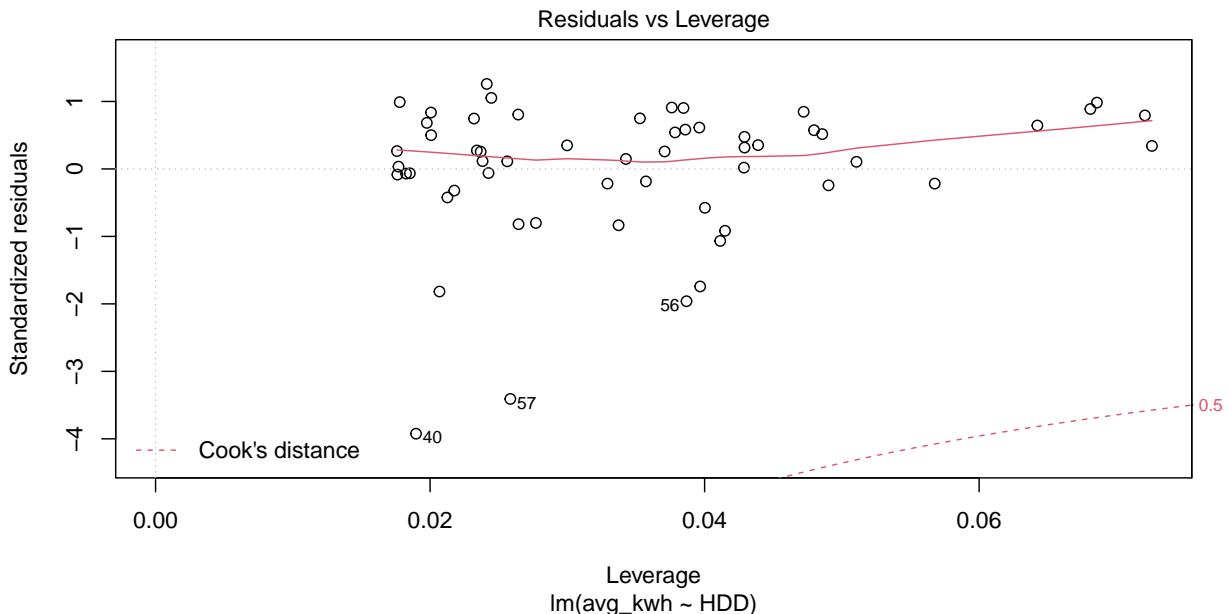
```
plot(avg_eff_HDD_auckland$HDD, avg_eff_HDD_auckland$avg_kwh)
abline(all_lm)
```



```
plot(all_lm)
```







```
all_mdl_lm = lm(data = eff_HDD_auckland, kwh~HDD+model )
summary(all_mdl_lm)
```

```

## Call:
## lm(formula = kwh ~ HDD + model, data = eff_HDD_auckland)
##
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -236.04   -72.33  -24.98   38.44 1222.80 
##
## Coefficients:
##                               Estimate Std. Error t value Pr(>|t|)    
## (Intercept)               131.930227   3.365507 39.201 < 2e-16 ***
## HDD                      0.012631    0.021162  0.597 0.550618  
## modelNissan Leaf (30 kWh) 55.081258   3.813767 14.443 < 2e-16 ***
## modelNissan Leaf (24 kWh) 2011-2012  3.595591   4.797043  0.750 0.453560  
## modelNissan Leaf (40 kWh) 32.951379   9.698679  3.398 0.000685 *** 
## modelNissan e-NV200 (24 kWh) 14.578854   7.571009  1.926 0.054201 .  
## modelHyundai Ioniq (EV)    13.273285  19.091074  0.695 0.486919  
## modelBMW i3                -31.437710  13.866842 -2.267 0.023420 *  
## modelHyundai Kona (EV)     -6.649758  15.782093 -0.421 0.673516  
## modelRenault Zoe           13.748701  18.217906  0.755 0.450471  
## modelTesla Model 3          53.107347  19.331772  2.747 0.006030 ** 
## modelNissan Leaf (62 kWh)  0.007987  53.615774  0.000 0.999881  
## modelKia Niro (EV)         -49.965591  21.308532 -2.345 0.019067 *  
## modelTesla Model S          98.177435  13.782264  7.123 1.18e-12 *** 
## modelTesla Model-X          78.572046  22.369796  3.512 0.000447 *** 
## modelKia Soul               114.152473  16.332153  6.989 3.07e-12 *** 
## modelMG ZS EV                -25.499294  37.940987 -0.672 0.501561  
## modelToyota Prius            -62.528363  14.834752 -4.215 2.54e-05 *** 

```

```

## modelMitsubishi Outlander          29.549372 69.170396  0.427 0.669252
## modelAudi A3 e-tron             86.121902 42.401335  2.031 0.042289 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 119.7 on 5860 degrees of freedom
## Multiple R-squared:  0.05966,   Adjusted R-squared:  0.05661
## F-statistic: 19.57 on 19 and 5860 DF,  p-value: < 2.2e-16

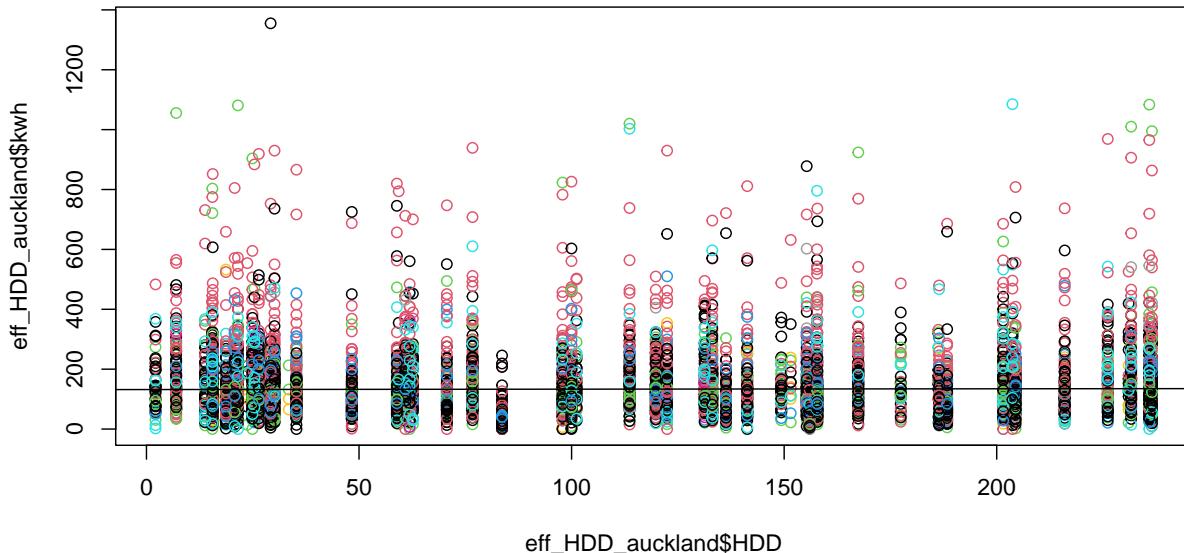
plot(eff_HDD_auckland$HDD, eff_HDD_auckland$kwh, col = eff_HDD_auckland$model)
abline(all_mdl_lm)

```

```

## Warning in abline(all_mdl_lm): only using the first two of 20 regression
## coefficients

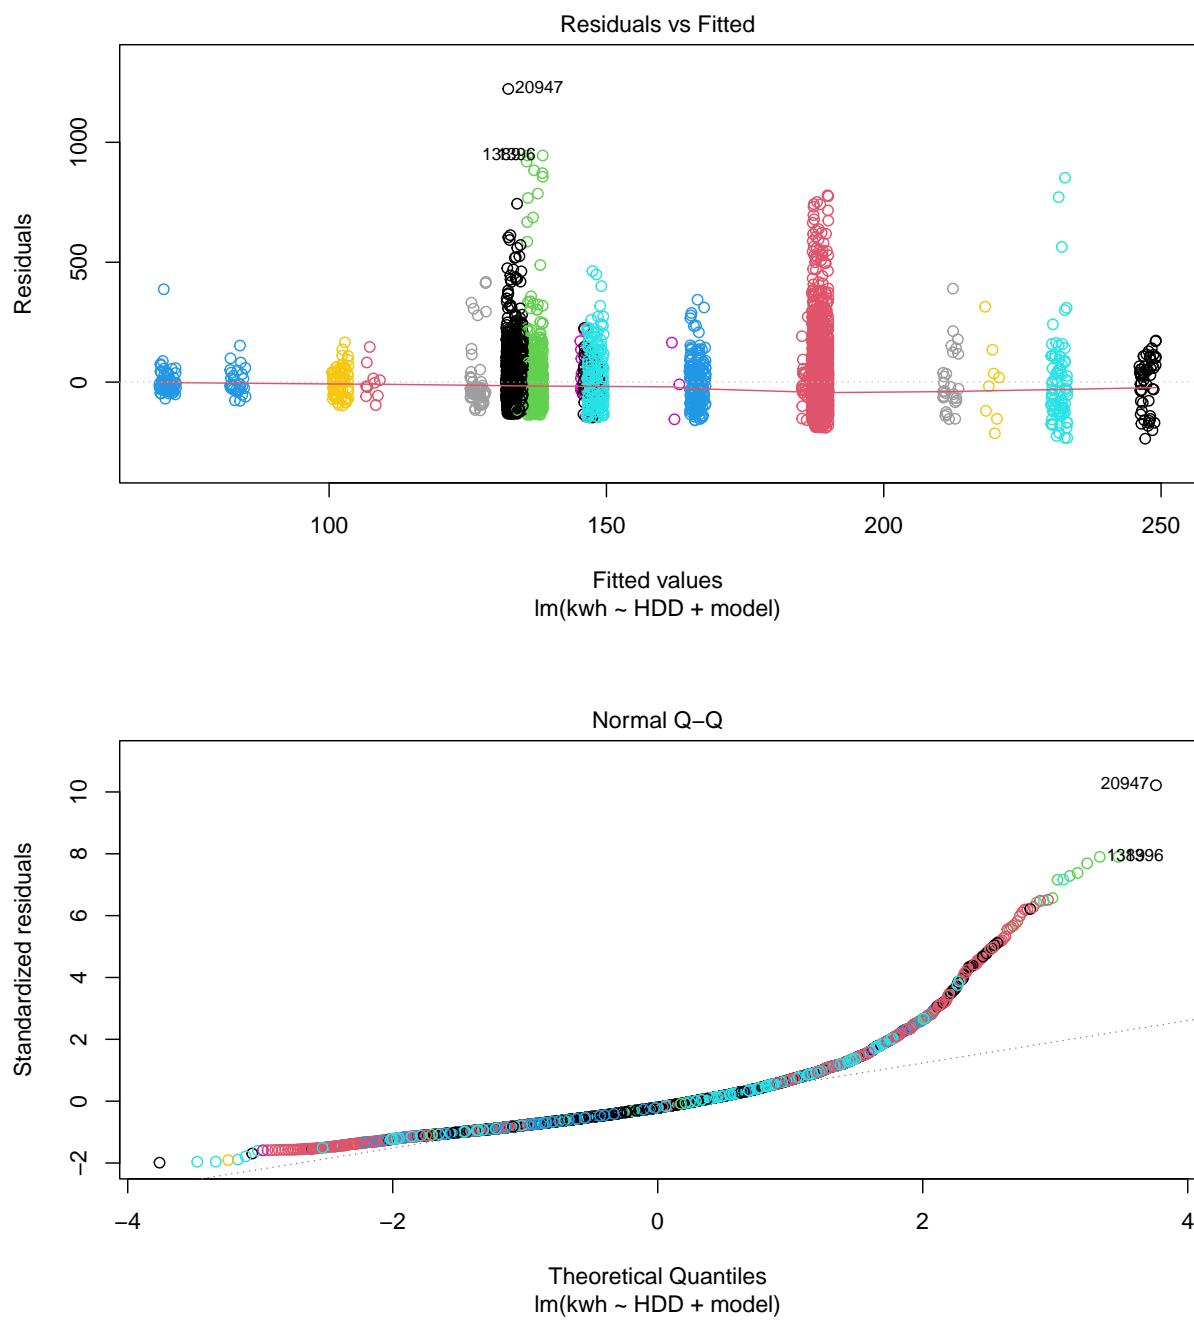
```

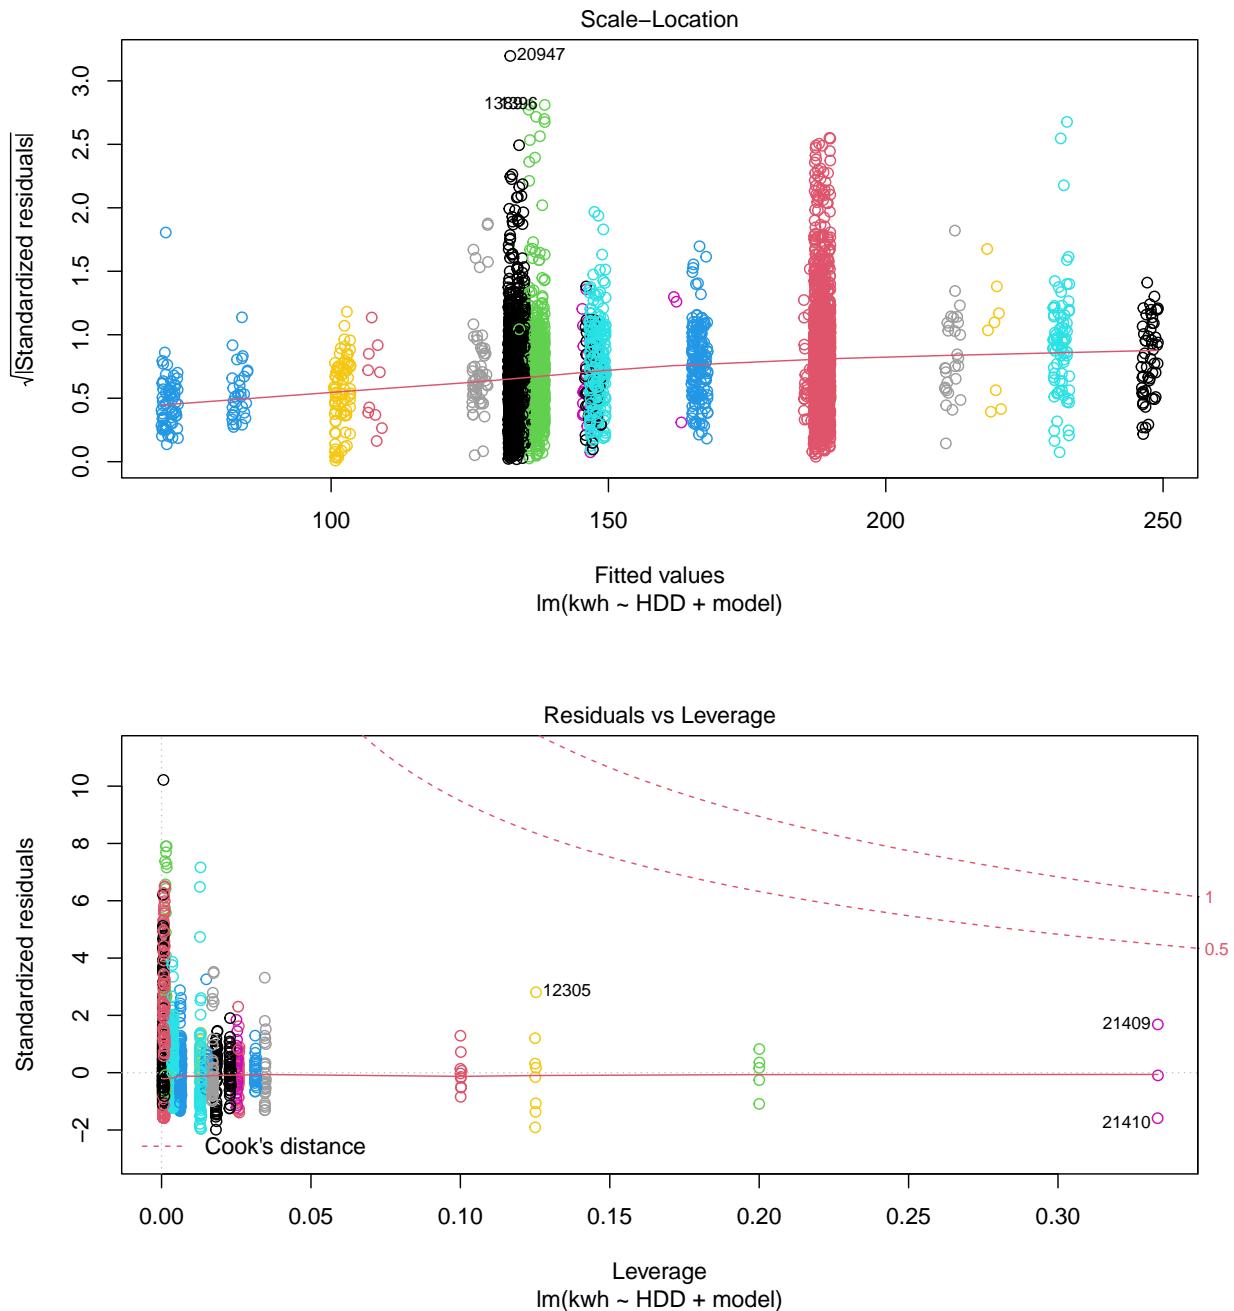


```

plot(all_mdl_lm, col = eff_HDD_auckland$model)

```





linear model to try to model kWh used by number of HDD in the month is effectively useless only thing that can be seen is the cars that are driven more or less each month

```
card_fuel = read.csv("Downloaded stats/fuel_usage_card_transactions_data.csv", header = T)
head(card_fuel)
```

```
##   year month fuel_purchased
## 1 2003     1      207.2
## 2 2003     2      206.2
## 3 2003     3      223.5
```

```

## 4 2003      4          198.3
## 5 2003      5          192.6
## 6 2003      6          180.9

fuel_cost = read.csv("Downloaded stats/MBIE_weekly_fuel_cost.csv")
fuel_cost$Week_ending_Friday = as.Date(fuel_cost$Week_ending_Friday)
fuel_cost$year = as.numeric(format(fuel_cost$Week_ending_Friday, '%Y'))
fuel_cost$month = as.numeric(format(fuel_cost$Week_ending_Friday, '%m'))
head(fuel_cost)

##   Week_ending_Friday Week_number Status Dubai_crude_USD.p.bbl NZ_US_exrate
## 1      2004-04-23       2004w17 Final        32.10      0.633
## 2      2004-04-30       2004w18 Final        32.65      0.626
## 3      2004-05-07       2004w19 Final        33.38      0.630
## 4      2004-05-14       2004w20 Final        34.41      0.609
## 5      2004-05-21       2004w21 Final        35.77      0.602
## 6      2004-05-28       2004w22 Final        35.03      0.618
##   Dubai_crude_NZD.p.bbl Diesel_importer_cost_NZc.p.l Diesel_ETS_NZc.p.l
## 1           50.74                  43.4                 0
## 2           52.16                  45.8                 0
## 3           52.94                  47.8                 0
## 4           56.55                  48.1                 0
## 5           59.40                  46.5                 0
## 6           56.68                  45.7                 0
##   Diesel_GST_NZc.p.l Diesel_taxes_NZc.p.l Diesel_price_excl_taxes_NZc.p.l
## 1             7.5                  7.9                 59.6
## 2             7.3                  7.7                 58.3
## 3             7.8                  8.2                 62.2
## 4             8.3                  8.7                 66.2
## 5             8.3                  8.6                 65.7
## 6             8.5                  8.9                 68.0
##   Diesel_main_port_price_NZc.p.l Diesel_discounted_retail_price_NZc.p.l
## 1                   67.5                 67.5
## 2                   66.0                 66.0
## 3                   70.3                 70.3
## 4                   74.9                 74.9
## 5                   74.3                 74.3
## 6                   76.9                 76.9
##   Diesel_importer_margin_NZc.p.l Diesel_margin_trend_NZc.p.l
## 1                   18.7                18.33931
## 2                   15.9                18.34124
## 3                   18.8                18.34317
## 4                   20.5                18.34510
## 5                   17.9                18.34703
## 6                   19.9                18.34897
##   Regular_Petrol_importer_cost_NZc.p.l Regular_Petrol_ETS_NZc.p.l
## 1                   48.6                  0
## 2                   52.2                  0
## 3                   53.3                  0
## 4                   56.4                  0
## 5                   54.3                  0
## 6                   53.7                  0
##   Regular_Petrol_GST_NZc.p.l Regular_Petrol_taxes_NZc.p.l
## 1                   12.7                54.7

```

## 2	12.6	54.5
## 3	13.0	55.0
## 4	13.5	55.5
## 5	13.5	55.4
## 6	13.8	55.7
## Regular_Petrol_price_excl_taxes_NZc.p.1		
## 1	59.8	
## 2	58.5	
## 3	62.3	
## 4	66.4	
## 5	65.9	
## 6	68.2	
## Regular_Petrol_main_port_price_NZc.p.1		
## 1	114.5	
## 2	113.0	
## 3	117.3	
## 4	121.9	
## 5	121.3	
## 6	123.9	
## Regular_Petrol_discounted_retail_price_NZc.p.1		
## 1	114.5	
## 2	113.0	
## 3	117.3	
## 4	121.9	
## 5	121.3	
## 6	123.9	
## Regular_Petrol_importer_margin_NZc.p.1	Regular_Petrol_margin_trend_NZc.p.1	
## 1	13.9	13.80219
## 2	11.8	13.78385
## 3	13.8	13.76550
## 4	14.2	13.74716
## 5	12.6	13.72881
## 6	11.8	13.71047
## Premium_Petrol_95R_importer_cost_NZc.p.1	Premium_Petrol_95R_ETS_NZc.p.1	
## 1	50.54357	0
## 2	52.86514	0
## 3	55.80943	0
## 4	57.75229	0
## 5	60.77043	0
## 6	57.88643	0
## Premium_Petrol_95R_GST_NZc.p.1	Premium_Petrol_95R_taxes_NZc.p.1	
## 1	13.10000	55.06500
## 2	13.43333	55.39833
## 3	14.10000	56.06500
## 4	14.10000	56.06500
## 5	14.10000	56.06500
## 6	14.32222	56.28722
## Premium_Petrol_95R_price_excl_taxes_NZc.p.1		
## 1	62.83500	
## 2	65.50167	
## 3	70.83500	
## 4	70.83500	
## 5	70.83500	
## 6	72.61278	

```

##  Premium_Petrol_95R_main_port_price_NZc.p.1
## 1                      117.9
## 2                      120.9
## 3                      126.9
## 4                      126.9
## 5                      126.9
## 6                      128.9
##  Premium_Petrol_95R_discounted_retail_price_NZc.p.1
## 1                      117.9
## 2                      120.9
## 3                      126.9
## 4                      126.9
## 5                      126.9
## 6                      128.9
##  Premium_Petrol_95R_importer_margin_NZc.p.1
## 1                      12.29143
## 2                      12.63652
## 3                      15.02557
## 4                      13.08271
## 5                      10.06457
## 6                      14.72635
##  Premium_Petrol_95R_margin_trend_NZc.p.1
## 1                      13.92260
## 2                      13.91374
## 3                      13.90487
## 4                      13.89600
## 5                      13.88714
## 6                      13.87827
##  Diesel_main_port_price_excl_taxes_NZc.p.1 year month
## 1                      59.6 2004     4
## 2                      58.3 2004     4
## 3                      62.1 2004     5
## 4                      66.2 2004     5
## 5                      65.7 2004     5
## 6                      68.0 2004     5

month_fuel_cost = fuel_cost %>%
  group_by(year, month) %>%
  summarise(petrol = mean(Regular_Petrol_discounted_retail_price_NZc.p.1), diesel = mean(Diesel_discounted_retail_price_NZc.p.1),
            Premium = mean(Premium_Petrol_95R_discounted_retail_price_NZc.p.1))

## 'summarise()' has grouped output by 'year'. You can override using the '.groups' argument.

head(month_fuel_cost)

## # A tibble: 6 x 5
## # Groups:   year [1]
##   year month petrol diesel Premium
##   <dbl> <dbl>  <dbl>  <dbl>   <dbl>
## 1 2004     4    114.   66.8    119.
## 2 2004     5    121.   74.1    127.
## 3 2004     6    120.   72.9    124.
## 4 2004     7    115.   71      120.

```

```

## 5 2004     8   122.    78.6    127.
## 6 2004     9   120.    78.2    125.

card_fuel = merge(card_fuel, month_fuel_cost, by = c("year", "month"), sort = FALSE)
card_fuel = card_fuel[card_fuel$year != 2004,]
head(card_fuel)

##   year month fuel_purchased petrol diesel Premium
## 10 2005     1       279.9 114.875 76.475 121.400
## 11 2005     2       277.0 118.050 76.550 123.150
## 12 2005     3       321.5 121.525 83.250 127.025
## 13 2005     4       300.7 128.540 88.000 133.900
## 14 2005     5       290.2 123.900 82.425 128.150
## 15 2005     6       293.2 126.250 86.250 132.900

#plot(card_fuel$month[card_fuel$year == 2005],card_fuel$fuel_purchased[card_fuel$year == 2005],
#      type = 'b', col = 1 , ylim = c(min(card_fuel$fuel_purchased),max(card_fuel$fuel_purchased)))
#for (i in min(card_fuel$year):max(card_fuel$year)) {
#  points(card_fuel$month[card_fuel$year == i],card_fuel$fuel_purchased[card_fuel$year == i],
#         type = 'b', col = i+1-min(card_fuel$year) )
#
#}

```

effect of lockdown is very obvious so probably have to remove 2020-2021

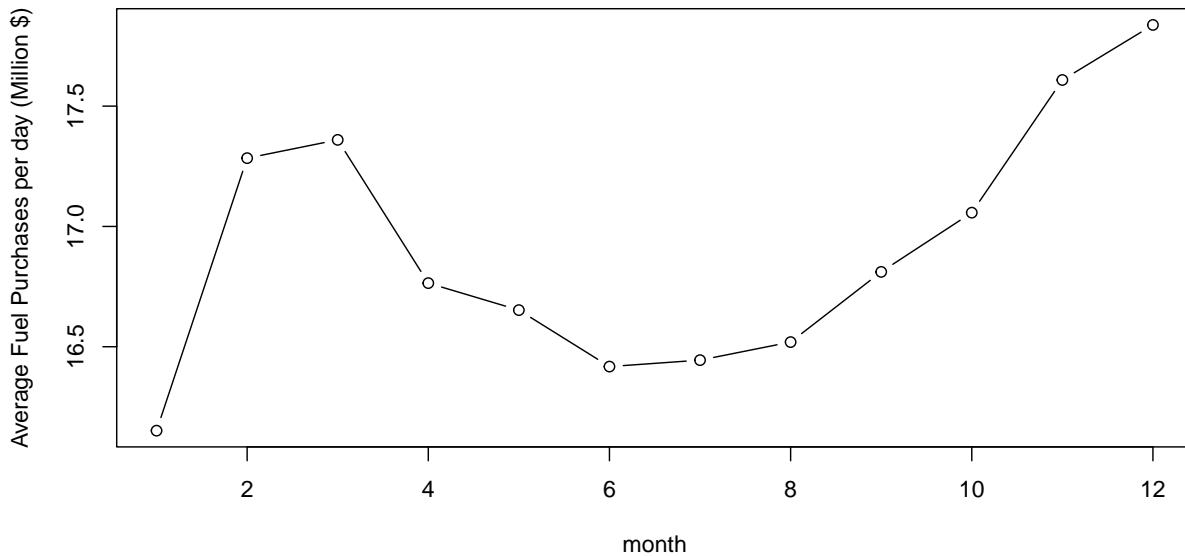
```

mean_card_fuel = card_fuel[card_fuel$year <=2019,] %>%
  group_by(month) %>%
  summarise(avg_fuel_purchased = mean(fuel_purchased))
head(mean_card_fuel)

## # A tibble: 6 x 2
##   month avg_fuel_purchased
##   <int>          <dbl>
## 1     1            501.
## 2     2            484.
## 3     3            538.
## 4     4            503.
## 5     5            516.
## 6     6            493.

plot(mean_card_fuel$month, mean_card_fuel$avg_fuel_purchased/month_length, type = 'b', xlab= 'month', y

```

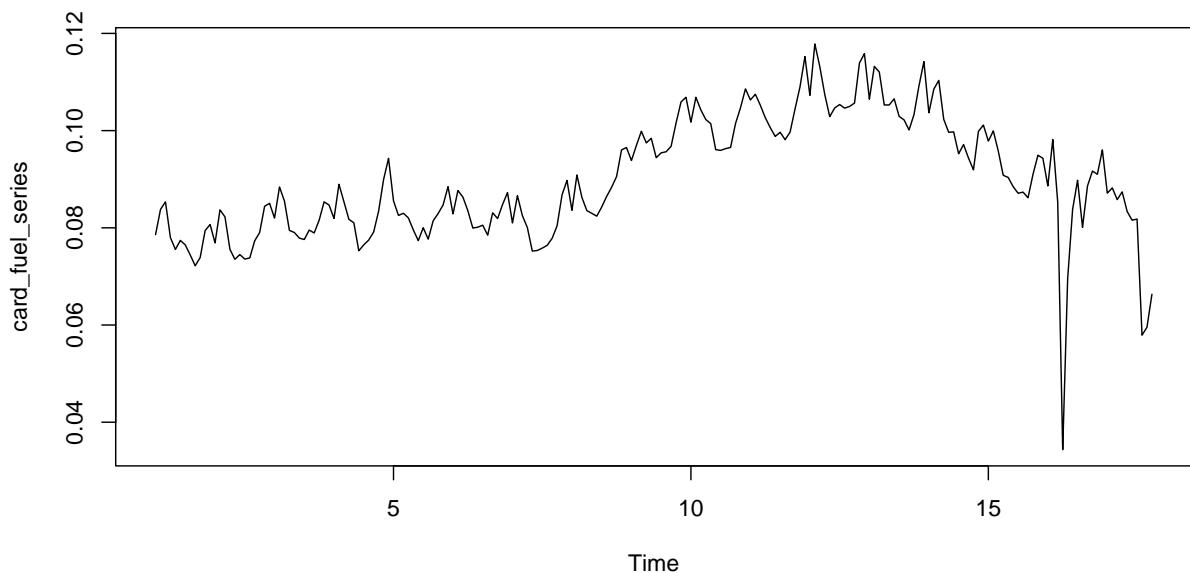


money spent shows a very clear upwards trend towards December (trips). with a valley in the winter. interesting that jan is lowest point. maybe new years resolutions to bike/walk or no school and not as many trips as dec?

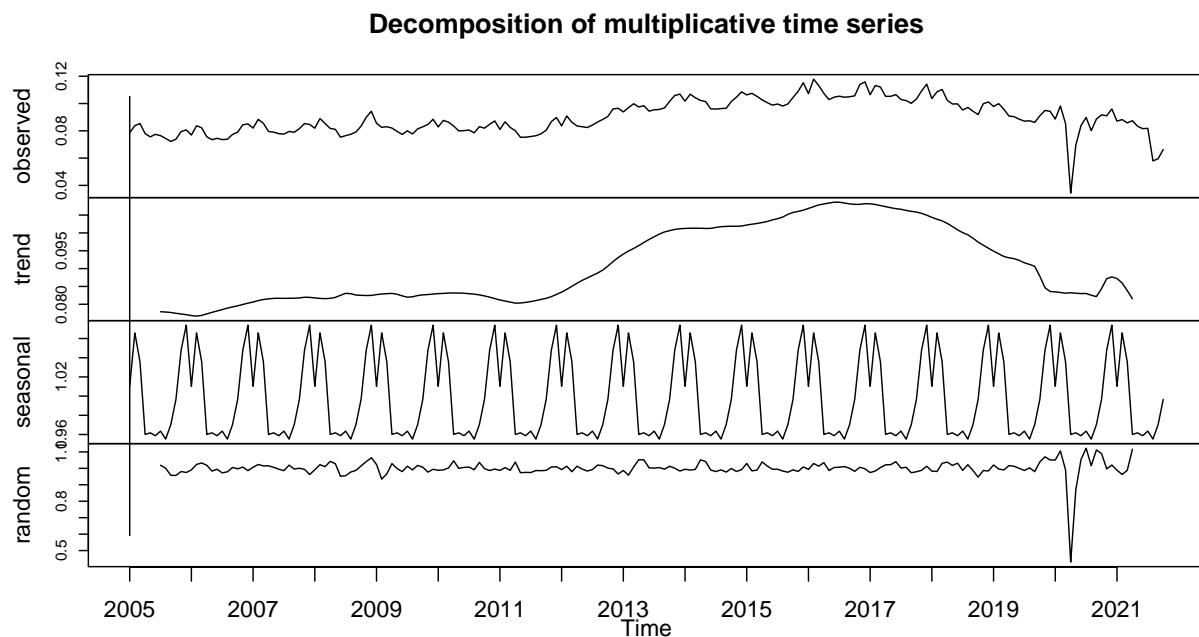
```
card_fuel_series = ts((card_fuel$fuel_purchased/card_fuel$petrol)/month_length, frequency = 12)

## Warning in (card_fuel$fuel_purchased/card_fuel$petrol)/month_length: longer
## object length is not a multiple of shorter object length

plot(card_fuel_series)
```



```
card_fuel_decomp = decompose(card_fuel_series, "multiplicative")
plot(card_fuel_decomp, xaxt = "n")
axis(1, labels = 2005:2021, at = 1:17, outer = T, pos = 0.520)
abline(v = 1, )
```



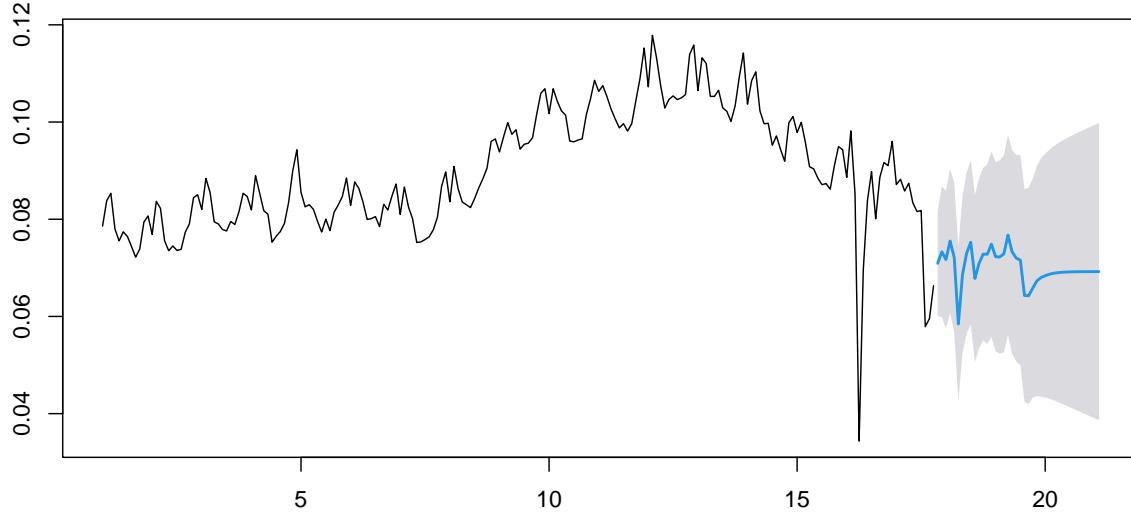
```
card_fuel_model = auto.arima(card_fuel_series) #uses AIC to get best model
summary(card_fuel_model)
```

```

## Series: card_fuel_series
## ARIMA(3,1,1)(0,0,2)[12]
##
## Coefficients:
##             ar1      ar2      ar3      ma1     sma1     sma2
##             0.6110 -0.1421  0.1089 -0.8782  0.1861  0.3320
## s.e.    0.0931  0.0830  0.0819   0.0583  0.0797  0.0986
## 
## sigma^2 estimated as 3.056e-05: log likelihood=760.91
## AIC=-1507.82  AICc=-1507.24  BIC=-1484.7
## 
## Training set error measures:
##               ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -0.0001161465 0.005431042 0.003179015 -0.7363641 4.08468 0.5848447
##                   ACF1
## Training set 0.004041763
## 
plot(forecast(card_fuel_model, level = 95, h = 40))

```

Forecasts from ARIMA(3,1,1)(0,0,2)[12]



```
plot(card_fuel_decomp$figure, type = '1')
```

